

UNIVERSITY OF PORTSMOUTH SCHOOL OF COMPUTING

Module Name:

M32363 Database Principle (DBPRIN) Lab Assignment (Individual)

Date:

2 Aug 2023

Intake: _

PT-UOP BScDSA-11_____

UOP ID Number: UP2200918

Contents

LAB 1 (16/07/2023) - Database Design.....	3
LAB 2 (16/07/2023) - Normalization	9
LAB 3 (26/07/2023) - Transactions	17
LAB 4 (19/07/2023) - Relational Algebra.....	23
LAB 5 (20/07/2023) - SQL, Query Optimization	29
LAB 6 (21/07/2023) - Optimization	37
LAB 8 (22/07/2023) -Data Manipulation and Query Output Formatting.....	43
LAB 9 (25/07/2023) - Advanced SQL.....	49
LAB 10 (25/07/2023) - Monitoring, Tuning, Backup & Recovery	59
LAB 11 (26/07/2023) - Christmas Fun.....	68

LAB 1 (16/07/2023) - Database Design

Lab Description

Lab 1 involves understanding and applying concepts related to the Entity-Relationship (ER) model, cardinality, and relationship types. Covers topics such as developing an ERD based on a scenario, identifying entity attributes and primary keys, determining cardinality and relationship symbols, and mapping ERD to table structures. The lab also includes SQL queries to retrieve specific information from the database based on the given ERD and table structure.

Q1. Which of the following steps occurs first in the process of building an ERD based on a scenario?





- a. Develop the initial ERD.
- b. Create a detailed narrative of the organisation's description of operations.
- c. Identify the attributes and primary keys that adequately describe the entities.
- d. Identify the business rules based on the description of operations.

Q2. A student can attend 5 modules. Different lecturers can offer the same module.

The relationship of students → lecturers is a _____ relationship.

- a. Many-to-many (M:M)
- b. One-to-many (1:M)
- c. One-to-one (1:1)
- d. Many-to-one (M:1)
- e. Many-to-zero OR one (M:0-1)
- f. One-to-one OR many (1:1-M)
- g. Many-to-zero

Q3. What would be the cardinality symbol for the relationship identified in Q2?

- a. 
- b. 
- c. 
- d. 

e. 

Q4. Which of the following statements best describes the function of an entity relation model?

- a. An ER model is concerned primarily with the view of the attributes that will be used in physical implementation
- b. An ER model is concerned primarily with a physical implementation of the data and secondly with the logical view
- c. An ER model provides a view of the logic of the data and not the physical implementation
- d. An ER model is entirely concerned with modelling the physical implementation

Q5. Consider Figure 1 representing instances of a relationship between EMPLOYEE and the DEPARTMENT that the employees work in.




Figure.1

Which of the following relationships are solved and represented by the above figure?

- a. **1:1** relationship between EMPLOYEE and DEPARTMENT and total participation from EMPLOYEE and total participation from DEPARTMENT
- b. **1:M** relationship between EMPLOYEE and DEPARTMENT and partial participation from EMPLOYEE and partial participation from DEPARTMENT
- c. **M:M** relationship between EMPLOYEE and DEPARTMENT and partial participation from EMPLOYEE and total participation from DEPARTMENT
- d. **M:1** relationship between EMPLOYEE and DEPARTMENT and total participation from EMPLOYEE and total participation from DEPARTMENT

Q6. Suppose we map the following ERD (Figure 2) to the relations T1(A, B) and T2(B, C). The CREATE table statement for T2 is defined as the following:

```
CREATE TABLE T2(B SERIAL PRIMARY KEY, C INT).
```




Figure. 2

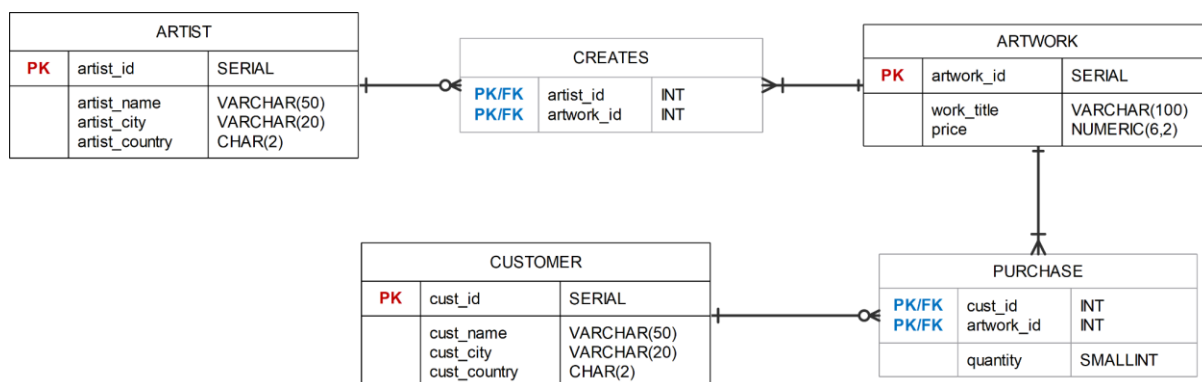
Which one of the following create table statements would be correct for T1, enforcing the FK and relationship?

- ```
CREATE TABLE T1(A SERIAL PRIMARY KEY UNIQUE, B INT, FOREIGN KEY NOT NULL);
```
- ```
CREATE TABLE T1(A SERIAL PRIMARY KEY, B INT NOT NULL REFERENCES T2(B));
```
- ```
CREATE TABLE T1(A SERIAL UNIQUE, B REFERENCES T2(B));
```
- ```
CREATE TABLE T1(A SERIAL, B INT);
```

Q7. In the manufacturing industry labourers are given different jobs on different days and each job has its own monthly basic and monthly special rates as wages to be paid to labourers. A worker is not given more than one type of job on a day. A database designer is given the job to design database for above situation and the designer designs a draft of one of the tables as :

FIELD	TYPE	REMARKS
From date	DATE	From this date to
To date	DATE	This date
Labour ID	SERIAL	The worker ID
Job done code	CHAR(6)	The job ID
At Basic Rate	DECIMAL(10,2)	At this basic rate
At Special Rate	DECIMAL(10,2)	At this special rate


Fig. 3



Draw an ERD for the above situation that represents the relationship between **job** and **labour** (place each attribute to the correct table along with data type/size).

Fig. 4

A:



Create a new database in your VM machine named **lab1** and paste the provided SQL code [[dbprin_lab1](#)] into the database. Analysing the provided ERD (Figure 4) and the code, provide answers for the following questions.

Note: You should format ALL your queries for the optimal output using appropriate column names and CONCAT where available (e.g. instead of “cus_name” it should be “Customer Name”).

Q8. Find the name of the artist who has created the Artwork titled '*Rainbow*'.

A:

[EDITABLE CODE]

```
SELECT a.artist_name
FROM artist a
JOIN creates c ON a.artist_id = c.artist_id
JOIN artwork art ON c.artwork_id = art.artwork_id
WHERE art.work_title = 'Rainbow';
```

[OUTPUT SCREENSHOT]

```
lab1=# SELECT a.artist_name
FROM artist a
JOIN creates c ON a.artist_id = c.artist_id
JOIN artwork art ON c.artwork_id = art.artwork_id
WHERE art.work_title = 'Rainbow';
 artist_name
-----
Ringo
(1 row)
```

Q9. Find the titles of the Artworks created by 'Lolo'.

A:

[EDITABLE CODE]

```
SELECT art.work_title
FROM artwork art
JOIN creates c ON art.artwork_id = c.artwork_id
JOIN artist a ON c.artist_id = a.artist_id
WHERE a.artist_name = 'Lolo';
```

[OUTPUT SCREENSHOT]

```

up2200918=# \c lab1
You are now connected to database "lab1" as user "up2200918".
lab1=# SELECT art.work_title
FROM artwork art
JOIN creates c ON art.artwork_id = c.artwork_id
JOIN artist a ON c.artist_id = a.artist_id
WHERE a.artist_name = 'Lolo';
       work_title
-----
 Colours of the Sky
 Long road to home
(2 rows)

```

Q10. Find the names of customers who have not bought an artwork priced more than £ 200. List the customer's name, artwork title and the price. Add the £ sign to the output and order on price from lowest to highest.

A:

[EDITABLE CODE]

```

SELECT c.cust_name, art.work_title, CONCAT('£', art.price) AS price
FROM customer c
JOIN purchase p ON c.cust_id = p.cust_id
JOIN artwork art ON p.artwork_id = art.artwork_id
WHERE art.price <= 200
ORDER BY art.price;

```

[OUTPUT SCREENSHOT]

```

lab1=# SELECT c.cust_name, art.work_title, CONCAT('£', art.price) AS price
FROM customer c
JOIN purchase p ON c.cust_id = p.cust_id
JOIN artwork art ON p.artwork_id = art.artwork_id
WHERE art.price <= 200
ORDER BY art.price;
 cust_name | work_title | price
-----
 Mary      | The war    | £1.00
 Michael   | Reflection | £40.00
 Mary      | Reflection | £40.00
 John      | Reflection | £40.00
 Sally     | The moon   | £145.00
 Margaret  | Night street | £150.00
 Michael   | Long road to home | £150.00
 Paul      | Long road to home | £150.00
 Sally     | Long road to home | £150.00
 Paul      | Night street | £150.00
 Harry     | My lovely song | £200.00
(11 rows)

```

Conclusion/Reflection

Lab 1 provided practical experience in database design using the ER model. I learned how to identify entities, attributes, and relationships in a scenario and represent them in an ERD. The lab also improved my skills in SQL query writing and database mapping. Overall, it gave me a strong foundation in database modeling and design.

LAB 2 (16/07/2023) - Normalization

Lab Description

Lab 2 involves normalizing tables to achieve 1NF and 3NF, creating an Entity-Relationship Diagram (ERD) based on the normalized tables, and working with data types and sizes. It also includes setting up a new database, creating tables with appropriate primary keys, foreign keys, constraints, and data types, as well as inserting sample data into the tables. The lab focuses on

order_id	order_date	cust_id	cust_name	cust_country	prod_id	prod_name	prod_price	prod_qty
1001 1001 1001	01/10/2022 01/10/2022 01/10/2022	1	Apple Apple Apple	US	7, 5, 4	table, desk, chair	800, 325, 200	1, 1, 5
1002 1002	02/10/2022 02/10/2022	2	Samsung	KO	11, 4	dresser, chair	500, 200	4, 2
1003	03/10/2022	3	Benq	DE	11	dresser	500	3

database normalisation principles and schema design.

Fig. 1

Q1. You have the following table (Fig.1) with data. Normalise the table in 1NF (show 1NF) and create the ERD for 3NF, with associated data type. You can use Visio to create diagram

Hint: *You should have 4 (or 5) tables for 3NF*

A:


order_id	order_date	cust_id	cust_name	cust_country	prod_id	prod_name	prod_price	prod_qty
1001	01/10/2022	1	Apple	US	7	table	800	1
1001	01/10/2022	1	Apple	US	5	desk	325	1
1001	01/10/2022	1	Apple	US	4	chair	200	5
1002	02/10/2022	2	Samsung	KO	11	dresser	500	4
1002	02/10/2022	2	Samsung	KO	4	chair	200	2
1003	03/10/2022	3	Benq	DE	11	dresser	500	3

PRODUCT (prod_id, prod_name, prod_price)

CUSTOMER (cust_id, cust_name, cust_country)

ORDER (order_id, cust_id*, order_date)

ORDER_PRODUCT (order_id*, prod_id*, prod_qty)



Q2. You are given the following table in 1NF. Normalise data in 3NF and create the ERD with all data type and size. One Employee works 1 or M projects and on each project a different time is allocated

dept_id	dep_name	emp_id	emp_name	project_id	project_name	budgeted_time
10	Finance	1	Harry	100	Alpha	4.5
10	Finance	5	Dewey	105	Beta	3
10	Finance	11	Louie	103	Gamma	7
20	R&D	2	Jack	107	Delta	8
20	R&D	4	Jill	102	Echo	9

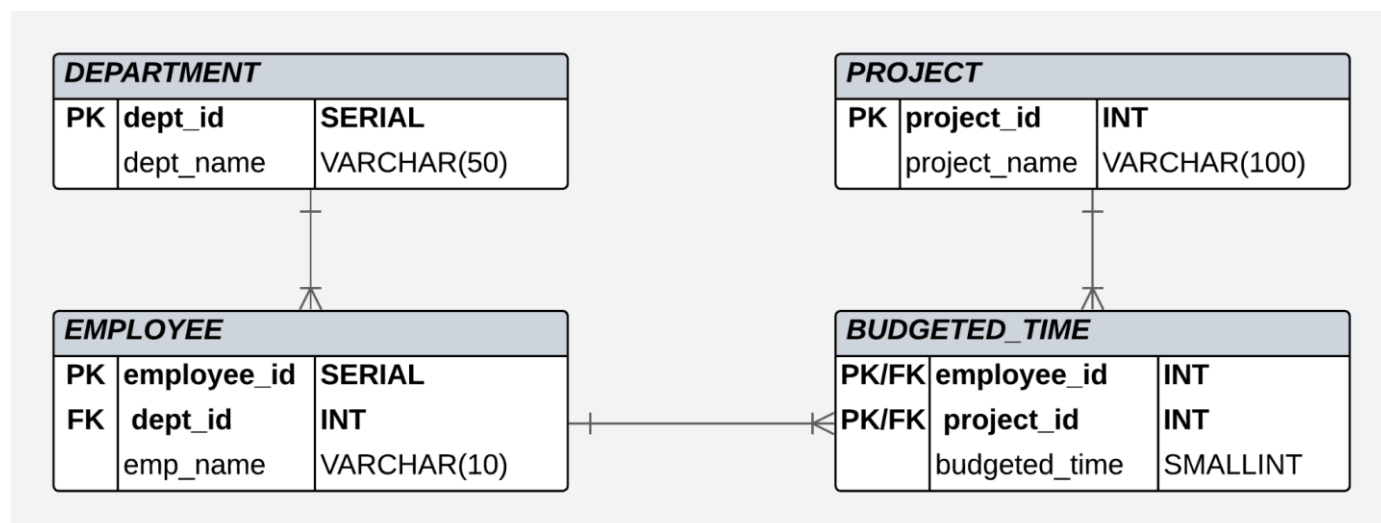
A:

DEPARTMENT(dept_id, dept_name)

EMPLOYEE(emp_id, dept_id, emp_name)

PROTECT(project_id, project_name)

BUDGET_TIME(emp_id*, project_id*, budgeted_time)



Q3. University of Portsmouth keeps the following details about a student and the various modules the student studied (not accurate):

up_number (student registration number to university)

stu_name (student name)

stu_addr (student address)

tut_id (tutor id)

tut_name (tutor name)

course_id - (course code)

course_name (course name)

module_id (module code)

module_name (module name)

module_results (module exam result)

Which of the following is one of the steps for normalizing **the relation Student to 1NF** as a student can have only one COURSE and one TUTOR but multiple modules?

Hint: you can input a quick sample output in a spreadsheet (like Q1) and remove the repeating group.

a. **STUDENT**(up_number, stu_name, stu_addr, tutor_id, tutor_name, course_id, course_name)
MODULE(module_id, module_name, module_results)

b. **STUDENT**(up_number, stu_name, stu_addr, tutor_id, tutor_name, course_id, course_name, module_code, module_results)

MODULE (module_id, module_name, course_id)

c. **STUDENT**(stu_name, stu_addr, tutor_id, tutor_name, course_id, module_id, module_name, module_results)

COURSE(up_number, course_id, course_name)

d. **STUDENT**(student_id, stu_name, stu_addr)

MODULE(tutor_id, module_id, module_name, module_results)

COURSE(course_id, course_name, tutor_id, tutor_name)

Q4. Based on the previous data sample what would be a 3NF normalisation? Write just the table name and the attributes as above.

A:

STUDENT(stu_id, stu_name, stu_addr)

MODULE(module_id, module_name, module_results)

COURSE(course_id, course_name)

TUTOR(tut_id, tut_name)

Q5. Open your VM and create a new database (lab2). You are given the following data dictionary.

ANIMAL_TYPE						
Attribute Name	Data Type	Size	Key	Reference	Constraints	Description
animal_type_id	SERIAL		PK			
common_name	VARCHAR	50			NOT NULL, UNIQUE	(eg. 'Arctic Wolf')
scientific_name	VARCHAR	150			NOT NULL	Official scientific name of the animal
conservation_status	VARCHAR	50			NOT NULL	('Endangered', 'Least Concerned')

MENAGERIE						
Attribute Name	Data Type	Size	Key	Reference	Constraints	Description
menagerie_id	SERIAL		PK			
common_name	VARCHAR	50	FK	animal_type > common_name	NOT NULL	
date_aquired	DATE				NOT NULL	
gender	CHAR	1			NOT NULL	
aquired_from	VARCHAR	250			NOT NULL	
name	VARCHAR	50			NOT NULL	
notes	TEXT					

menagerie_id	common_name	date_acquired	gender	acquired_from	name	notes
1	Bengal Tiger	2011-07-14	F	Dhaka Zoo	Ariel	Healthy coat at last exam.
2	Arctic Wolf	2008-09-30	F	National Zoo	Freddy	Strong appetite.
3	Bengal Tiger	2006-06-01	M	Scotland Zoo	Spark	Likes to play
4	Arctic Wolf	2007-06-12	F	Southampton National Park	Mia	Doesn't like sun

and the data sample output:

- Write the CREATE statements for tables ANIMAL_TYPE and MENAGERIE, including PKs, FKs, constraints, data type and size.

A:

[EDITABLE CODE]

```
CREATE TABLE ANIMAL_TYPE (
    animal_type_id SERIAL NOT NULL PRIMARY KEY,
    common_name VARCHAR(50) NOT NULL UNIQUE,
    scientific_name VARCHAR(150) NOT NULL,
    conservation_status VARCHAR(50) NOT NULL
);

CREATE TABLE MENAGERIE (
    menagerie_id SERIAL NOT NULL PRIMARY KEY,
    common_name VARCHAR(50) NOT NULL,
    FOREIGN KEY(common_name) REFERENCES ANIMAL_TYPE (common_name),
    date_acquired DATE NOT NULL,
    gender CHAR(1) NOT NULL,
    acquired_from VARCHAR(250) NOT NULL,
    name VARCHAR(50) NOT NULL,
    notes TEXT
);
```

[OUTPUT SCREENSHOT]

```
lab2=# CREATE TABLE ANIMAL_TYPE (
    animal_type_id SERIAL NOT NULL PRIMARY KEY,
    common_name VARCHAR(50) NOT NULL UNIQUE,
    scientific_name VARCHAR(150) NOT NULL,
    conservation_status VARCHAR(50) NOT NULL
);
CREATE TABLE
lab2=# CREATE TABLE MENAGERIE (
    menagerie_id SERIAL NOT NULL PRIMARY KEY,
    common_name VARCHAR(50) NOT NULL,
    FOREIGN KEY(common_name) REFERENCES ANIMAL_TYPE (common_name),
    date_acquired DATE NOT NULL,
    gender CHAR(1) NOT NULL,
    acquired_from VARCHAR(250) NOT NULL,
    name VARCHAR(50) NOT NULL,
    notes TEXT
);
CREATE TABLE
```

- Write 6 INSERT STATEMENTS for the following:
 - Common Name: 'Bengal Tiger', 'Arctic Wolf'
 - Scientific Name: 'Panthera tigris tigris', 'Canis lupus arctos'
 - Conservation Status: 'Endangered', 'Least Concern'
 - Acquired Date: '14/07/2011', '30/09/2008', '01/06/2006', '12/06/2007'
 - Gender: 'M', 'F'
 - Acquired From: 'Dhaka Zoo', 'National Zoo', 'Scotland Zoo', 'Southampton National Park'
 - Name: 'Ariel', 'Freddy', 'Spark', 'Mia'
 - Notes: 'Healthy coat at last exam', 'Strong appetite', 'Likes to play', 'Doesn't like sun'

The output of both tables should be exactly like in provided examples for ANIMAL_TYPE and MENAGERIE.

A:

[EDITABLE CODE]

```
INSERT INTO ANIMAL_TYPE
VALUES
(1, 'Bengal Tiger', 'Panthera tigris tigris', 'Endangered'),
(2, 'Arctic Wolf', 'Canis lupus arctos', 'Least Concern');

INSERT INTO MENAGERIE
VALUES
(1, 'Bengal Tiger', '2011-07-14', 'F', 'Dhaka Zoo', 'Ariel', 'Healthy coat at last exam'),
(2, 'Arctic Wolf', '2008-09-30', 'F', 'National Zoo', 'Freddy', 'Strong appetite'),
(3, 'Bengal Tiger', '2006-06-01', 'M', 'Scotland Zoo', 'Spark', 'Likes to play'),
(4, 'Arctic Wolf', '2011-06-12', 'F', 'Southampton National Park', 'Mia', 'Doesn't like sun');
```

[OUTPUT SCREENSHOT]

```

lab2=# INSERT INTO ANIMAL_TYPE
VALUES
(1, 'Bengal Tiger', 'Panthera tigris tigris', 'Endangered'),
(2, 'Arctic Wolf', 'Canis lupus arctos', 'Least Concern');

INSERT INTO MENAGERIE
VALUES
(1, 'Bengal Tiger', '2011-07-14', 'F', 'Dhaka Zoo', 'Ariel', 'Healthy coat at last exam'),
(2, 'Arctic Wolf', '2008-09-30', 'F', 'National Zoo', 'Freddy', 'Strong appetite'),
(3, 'Bengal Tiger', '2006-06-01', 'M', 'Scotland Zoo', 'Spark', 'Likes to play'),
(4, 'Arctic Wolf', '2011-06-12', 'F', 'Southampton National Park', 'Mia', 'Doesn't like sun');

INSERT 0 2
INSERT 0 4
lab2=# SELECT * FROM animal_type;
 animal_type_id | common_name | scientific_name | conservation_status
-----+-----+-----+-----
          1 | Bengal Tiger | Panthera tigris tigris | Endangered
          2 | Arctic Wolf  | Canis lupus arctos   | Least Concern
(2 rows)

lab2=# SELECT * FROM MENAGERIE;
 menagerie_id | common_name | date_acquired | gender | acquired_from | name | notes
-----+-----+-----+-----+-----+-----+-----
          1 | Bengal Tiger | 2011-07-14    | F      | Dhaka Zoo     | Ariel | Healthy coat at last exam
          2 | Arctic Wolf  | 2008-09-30    | F      | National Zoo  | Freddy | Strong appetite
          3 | Bengal Tiger | 2006-06-01    | M      | Scotland Zoo  | Spark | Likes to play
          4 | Arctic Wolf  | 2011-06-12    | F      | Southampton National Park | Mia | Doesn't like sun
(4 rows)

```

Q6. Based on the previous database you have created, list all the animals that are endangered, along with common name, scientific name, animal name and date acquired.

A:

[EDITABLE CODE]

```

SELECT
    a.common_name AS "Common Name",
    a.scientific_name AS "Scientific Name",
    m."name" AS "Animal Name",
    m.date_acquired AS "Date Acquired"
FROM animal_type a
INNER JOIN menagerie m ON a.common_name = m.common_name
WHERE
    a.conservation_status = 'Endangered';

```

[OUTPUT SCREENSHOT]

```

lab2=# SELECT
    a.common_name AS "Common Name",
    a.scientific_name AS "Scientific Name",
    m."name" AS "Animal Name",
    m.date_acquired AS "Date Acquired"
FROM
    animal_type a
    INNER JOIN menagerie m on a.common_name = m.common_name
WHERE
    a.conservation_status = 'Endangered';
 Common Name | Scientific Name | Animal Name | Date Acquired
-----+-----+-----+-----
 Bengal Tiger | Panthera tigris tigris | Ariel      | 2011-07-14
 Bengal Tiger | Panthera tigris tigris | Spark      | 2006-06-01
(2 rows)

```

Conclusion/Reflection

Lab 2 provided hands-on experience in database normalization and schema design. I learned how to eliminate data redundancy and organize data efficiently. Creating an ERD helped me visualize relationships between entities. Setting up the database and writing SQL queries improved my practical skills in database implementation.

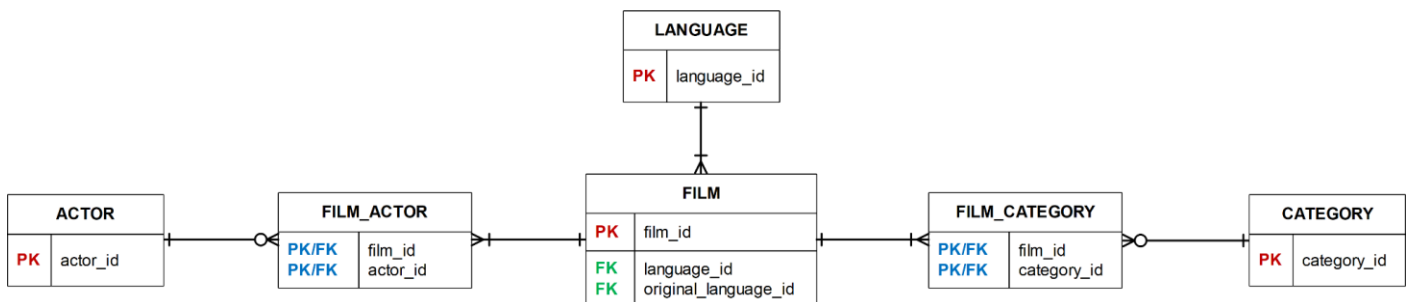
LAB 3 (26/07/2023) - Transactions

Lab Description

Lab 3 focused on database transactions and normalization. Involved building a database based on an ERD and performing various operations using SQL queries such as SELECT, INSERT, UPDATE, and DELETE. The lab also included the creation of a Transaction Analysis Matrix to analyze transactional behavior and its impact on the database.

- Q1.** You are given the following ERD without any additional attributes. Build up your database using your VM and **write 4 separate queries** that will contain a **SELECT**, **INSERT**, **UPDATE** and **DELETE** that will query the database. Create the **Transaction Analysis Matrix** for those transactions.

Do not spend time making the matrix “pretty” but accurate. You can use a spreadsheet or a table inserted in a document. Examples of the matrix are in the lecture presentation. Don't forget to add your work to the logbook.



(Query example - You might assume that some attributes are there (e.g. actor_name; actor_lname etc)

List the name of all movies where the actor's id (or surname) is X.

```

SELECT
    actor_name, film_name
FROM
    table_name
JOIN table_name ON table1.attr = table2.attr
JOIN table_name ON table3.attr = table4.attr
  
```

WHERE

table.attr = 'X';

Same for an new INSERT and for an UPDATE

A:

[EDITABLE CODE]

```
-- SELECT
SELECT f.film_id, f.title
FROM film f
JOIN film_actor fa ON f.film_id = fa.film_id
JOIN actor a ON fa.actor_id = a.actor_id
WHERE a.last_name = 'X';

-- INSERT
INSERT INTO actor (actor_id, first_name, last_name)
VALUES (1, 'John', 'Doe');

-- UPDATE
UPDATE actor
SET last_name = 'Smith'
WHERE actor_id = 1;

-- DELETE
DELETE FROM actor
WHERE actor_id = 1;
```

Transaction Analysis Matrix:

	actor	film_actor	film
SELECT	READ (a)	READ (fa)	READ (f)
INSERT	WRITE (a)	--	--
UPDATE	READ/WRITE (a)	--	--
DELETE	READ/WRITE (a)	--	--

Q2. Define in your own understanding (do not copy/paste definitions) of these terms: atomicity, consistency, isolation, durability, schedule, blind write, dirty read, unrepeatable read, serializable schedule, recoverable schedule, avoids-cascading-aborts schedule.

Note: *Some of the answers were in lecture, some in the book.*

A:

- **Atomicity:** Ensures transactions are all-or-nothing.
- **Consistency:** Keeps data valid and according to predefined rules.
- **Isolation:** Makes each transaction operate independently of others.
- **Durability:** Makes committed transactions survive system failures.
- **Schedule:** Defines the order of transaction execution.
- **Blind write:** An update operation that doesn't require reading the data first.
- **Dirty read:** Reading data from an uncommitted transaction.
- **Unrepeatable read:** Getting different values from reading the same data due to concurrent changes.
- **Serializable schedule:** Makes concurrent transactions equivalent to sequential ones.
- **Recoverable schedule:** Ensures database recovery after transaction failures.
- **Avoids-cascading-aborts schedule:** Prevents cascading rollbacks to preserve data consistency.

Q3. Consider the situation below, in which a number of account records have the following values:

$A1 = £40$; $A2 = £50$; $A3 = £30$

To transfer £10 from A3 to A1 while concurrently calculating the total funds in the three accounts, the following sequence of events may occur. Show the value of each data item in the last column, and discuss the reason for an incorrect summary value.

Note: Look at the lecture example and/or chapter 22 in Connolly and Begg.

Time	Transaction1	Transaction2	SUM
t ₁	sum = 0		
t ₂	read(A1)		
t ₃	sum = sum + A1		
t ₄		read(A3)	
t ₅		A3 = A3 - 10	
t ₆		write(A3)	
t ₇		read(A1)	
t ₈		A1 = A1 + 10	
t ₉		write(A1)	
t ₁₀		commit;	
t ₁₁	read(A3)		
t ₁₂	sum = sum + A3		
t ₁₃	read(A2)		
t ₁₄	sum = sum + A2		

A:

Time	Transaction1	Transaction2	A1	A2	A3	SUM
t ₁	sum = 0		40	50	30	0
t ₂	read(A1)		40	50	30	0
t ₃	sum = sum + A1		40	50	30	40
t ₄		read(A3)	40	50	30	40
t ₅		A3 = A3 - 10	40	50	30	40
t ₆		write(A3)	40	50	20	40
t ₇		read(A1)	40	50	20	40
t ₈		A1 = A1 + 10	40	50	20	50
t ₉		write(A1)	50	50	20	50
t ₁₀		commit;	50	50	20	50
t ₁₁	read(A3)		50	50	20	50
t ₁₂	sum = sum + A3		50	50	20	60
t ₁₃	read(A2)		50	50	20	60
t ₁₄	sum = sum + A2		50	50	20	110
t ₁₅	commit;					

Conclusion/Reflection

In Lab 3, I gained practical experience in handling transactions and implementing normalization techniques. I learned about transaction properties and how to ensure data consistency and integrity. The lab deepened my understanding of concurrency control and the importance of efficient database design. Overall, it enhanced my skills in database management and query execution.

LAB 4 (19/07/2023) - Relational Algebra

Lab Description

Lab 4 focused on SQL query formulation, its translation into relational algebra expressions, query tree drawings, and size estimations of query results, thereby enhancing understanding of database operations and data manipulation.

Q1. You are given the following query:

```
SELECT first_name, last_name, hire_date
FROM Employees
WHERE
    Title = 'Sales Representative';
```

Write the correspondent RA with π and σ and draw the tree diagram to show this relation

Relational Algebra Symbols:

σ SELECT (or RESTRICT) - Sigma

π PROJECT - Pi

\cup UNION

\cap INTERSECTION

- DIFFERENCE

\times PRODUCT (Cartesian Product)

\bowtie \Join \bowtie JOINS - Theta

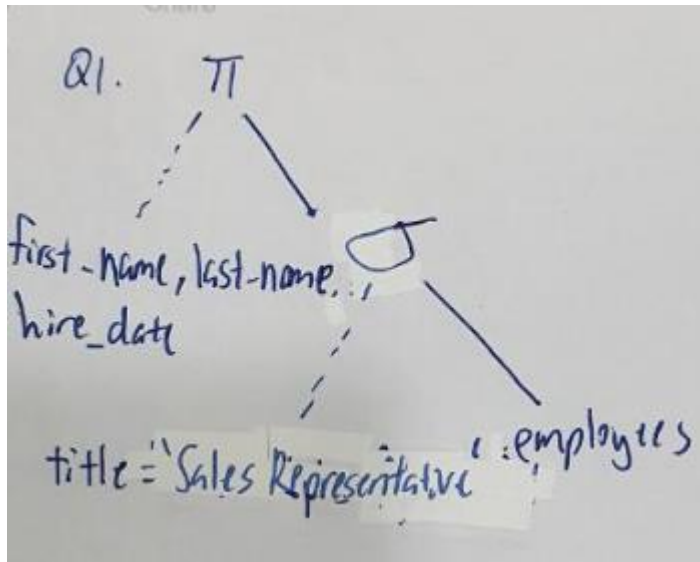
τ ORDER (SORT) - Tau

A:

Relational Algebra Expression:

$\pi(\text{first_name, last_name, hire_date})(\sigma_{\text{Title}='Sales Representative'}(\text{Employees}))$

Tree Diagram:



- Q2.** Let's expand the above query and we are adding a second WHERE clause filter.
Write the new RA and the tree diagram.

```

Select first_name, last_name,   hire_date
From Employees
Where
    Title = 'Sales Representative'
    and Country = 'UK';

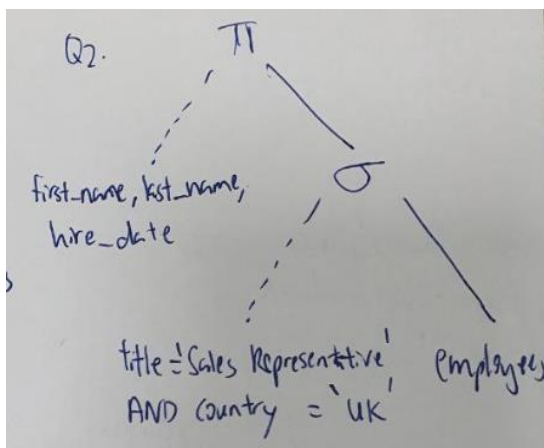
```

A:

Relational Algebra Expression:

$\pi(\text{first_name, last_name, hire_date})(\sigma_{\text{Title}='Sales Representative' \text{ and } \text{Country}='UK'}(\text{Employees}))$

Tree Diagram:



Q3. The following query has a conditional WHERE clause and ordering. What would be the RA for this?

```
SELECT order_id, order_date, shipper
FROM orders INNER JOIN shippers ON shippers . shipper_id = orders . ship_via
WHERE order_id < 10300
ORDER BY order_id
```


A:

Relational Algebra Expression:

τ (order_id)

π (order_id, order_date, shipper)

$(\sigma_{\text{order_id} < 10300} (\text{orders} \bowtie \text{shippers.shipper_id} = \text{orders.ship_via} \text{ shippers}))$



Q4. Consider a database with the following schema:

USER	(<u>name</u> , age, gender) PK - name
FREQUENTS	(<u>name</u> , <u>pizzeria</u>) PK name,pizzeria
EATS	(<u>name</u> , <u>pizza</u>) PK - name, pizza
SERVES	(<u>pizzeria</u> , <u>pizza</u> , price) PK - pizzeria, pizza

Write the **query**, **relational algebra expressions** and **query tree** for the following output.

“Find the names of all females who eat either mushroom or pepperoni pizza (or both).”

Hint: You can quickly draw the ERD using pen and paper for the above schema for an easier visualisation.

A:

Query:

```

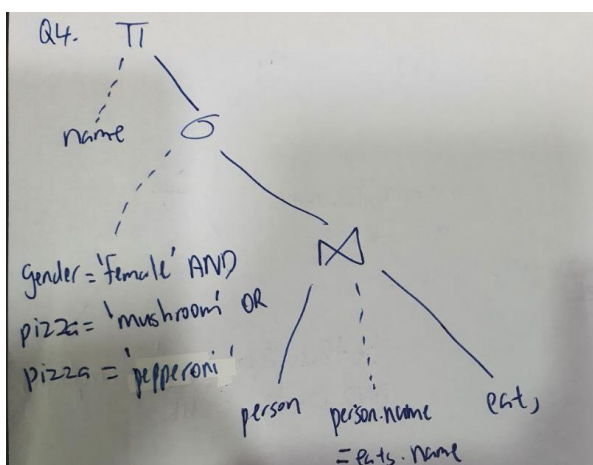
SELECT name
FROM user u
INNER JOIN eats e ON u.name = e.name
WHERE gender = 'female' AND
pizza = 'mushroom' OR pizza = 'pepperoni'

```

Relational Algebra Expressions:

$\pi(\text{name}) (\sigma_{\text{gender}='female' \ \& \ \text{pizza}='mushroom' \mid \text{pizza}='pepperoni'}}(\text{USER} \bowtie \text{EATS}))$

Query Tree:



Q5. Consider a database with the following schema:

PASSENGER (**p_id**, p_name, p_gender, p_city)
AGENCY (**a_id**, a_name, a_city)
FLIGHT (**f_id**, f_date, time, src, dest)
BOOKING (**p_id***, **a_id***, **f_id***, **f_date**)

Write the relational algebra expression and query tree for the following request:


“List details of all male passengers who are booked through operator”

A:

Relational Algebra Expressions:

$\pi(p_id, p_name, p_city)(\sigma_{p_gender='male'}(PASSENGER \bowtie BOOKING \bowtie AGENCY))$

Query Tree:



Q6. Relational algebra is a _____ Data Manipulation Language (DML).

- Declarative
- Object-oriented
- Procedural**
- Static
- None of the above

- Q7. Assuming that we have two relations R_1 and R_2 where R_1 has 10 tuples (records) and R_2 5, how many records will be there in the result of the operation $R_1 \times R_2$ (cartesian product)?
- a. 10 Records
 - b. 5 Records
 - c. 15 Records
 - d. 50 Records**
 - e. None

- Q8. For relation *Student*(*up, name, school*), suppose school can be one of {soc, smp, sces, smde, seee, dlw, icg}.

If the number of tuples in relation to Student is 1000, then what is the size estimate of the query

$\sigma_{\text{school} = \text{'physics'}}(\text{Student})?$

A:

0 tuples because Physics is not one of the schools. Hence, no students belong to that school.

Conclusion/Reflection

In this lab, I honed my skills in writing SQL queries and their equivalent relational algebra expressions, enhancing my database operation capabilities. Additionally, it provided insights into query optimization and size estimation. This practical application of theoretical concepts has improved my data retrieval and problem-solving skills in database management.

LAB 5 (20/07/2023) - SQL, Query Optimization

Lab Description

Lab 5 focused on practical SQL concepts, with the installation of a new database, Movie Rental, serving as a primary resource. This lab involved executing SQL queries on this database, performing operations like counting views, querying movie titles, and generating reports. Furthermore, we dove into the concept of query optimization, considering the performance of different queries. Lastly, hands-on tasks such as identifying and correcting SQL syntax errors were also included.

- Q1.** The below SQL statement has five syntax errors. Can you identify them? What would be the correct statement to return exactly the same result as in Fig1? Write the correct SQL statement.

```
SELECT
  movie_id AS "Movie ID",
  movie_title AS Movie_Title
  movie_lang AS 'LANGUAGE',
  cat_name AS "Category"
FROM
  movie
  INNER JOIN category LINK movie.movie_id = category.movie_id
WHERE
  movie_title = ALADDIN CALENDAR
```

Movie ID	movie_title	language	Category
10	ALADDIN CALENDAR	English	Sports
(1 row)			

Fig1.

A:

Syntax errors:

1. Missing commas between column aliases:

SQL requires commas to separate each column in the SELECT clause. In your query, you missed a comma between movie_title AS Movie_Title and movie_lang AS 'LANGUAGE'.

2. Incorrect use of 'LINK' instead of 'ON' in the INNER JOIN syntax:

The correct keyword to use when specifying the join condition is 'ON', not 'LINK'. So, your query should have INNER JOIN category ON movie.movie_id = category.movie_id instead of INNER JOIN category LINK movie.movie_id = category.movie_id.

3. Incorrect column name movie_lang:

There is no movie_lang column in the movie table, this will result in a syntax error. You need to replace movie_lang with the correct column name (name) from your language table that stores the language information.

4. Missing single quotes for the string in the WHERE clause:

In SQL, string values need to be enclosed in single quotes. So, you should have movie_title = 'ALADDIN CALENDAR' instead of movie_title = ALADDIN CALENDAR.

5. Missing semicolon at the end of the statement:

SQL statements should end with a semicolon. So, you should have **movie_title = 'ALADDIN CALENDAR';** instead of **movie_title = ALADDIN CALENDAR.**

[EDITABLE CODE]

```
SELECT
  m.movie_id AS "Movie ID",
  m.title AS "movie_title",
  l.name AS "language",
  cat.name AS "Category"
FROM
  movie m
INNER JOIN language l on l.language_id = m.language_id
INNER JOIN movie_Category mc on m.movie_id = mc.movie_id
INNER JOIN category cat ON mc.category_id = cat.category_id
WHERE
  m.title = 'ALADDIN CALENDAR';
```

[OUTPUT SCREENSHOT]

```

movie_rental=# SELECT
  m.movie_id AS "Movie ID",
  m.title AS "movie_title",
  l.name AS "language",
  cat.name AS "Category"
FROM
  movie m
INNER JOIN language l on l.language_id = m.language_id
INNER JOIN movie_Category mc on m.movie_id = mc.movie_id
INNER JOIN category cat ON mc.category_id = cat.category_id
WHERE
  m.title = 'ALADDIN CALENDAR';
Movie ID | movie_title | language | Category
-----+-----+-----+-----
      10 | ALADDIN CALENDAR | English | Sports
(1 row)

```

- Q2.** You are given the following two queries. The table employees have numerous attributes and your department needs the name, NIN and hire date of an employee. Which one will have a better performance and why?

<p>A</p> <pre> SELECT * FROM employee WHERE employee_id = 1; </pre>	<p>B</p> <pre> SELECT emp_name, emp_last_name, emp_nin, emp_hire_date from employee WHERE employee_id = 1; </pre>
--	--

A:

Query B performs better because it only retrieves specific columns (name, last_name, nin, hire_date), while Query A retrieves all columns with SELECT *. Fetching less data results in faster performance.

From this point forward we will install and use a new database (Movie Rental). This database will be your main DB for the rest of the module hence it is mandatory to install it for this and future labs.

Please follow [these instructions](#) on how to install the database. [The full ERD](#) (or you can find it on Moodle) and an [Interactive Data Dictionary](#).



DO NOT COPY/PASTE THE SQL STATEMENT AS WILL TAKE OVER 2 HOURS TO INSTALL IT (... yes I have counted). USE SERVER SIDE COMMAND LINE(S) TO INSTALL IN 20-30 SECONDS.

Q3. Using DB specific commands, How many **VIEWS** are inside of the movie_rental database?

A:

7 VIEWS

[EDITABLE CODE]

```
SELECT count(*) FROM information_schema.views WHERE table_schema = 'public';
```

[OUTPUT SCREENSHOT]

```
movie_rental=# SELECT count(*) FROM information_schema.views WHERE table_schema = 'public';
count
-----
      7
(1 row)
```

Q4. Using the provided ERD and Data Dictionary, create a query that will return all the movie titles that are in the **Documentary** category. Create 2 versions of the same query and in one use category as string ('documentary') and in the other use category ID (you will need to check what is the category ID for **Documentary**). Use the **EXPLAIN ANALYZE** command and see which query performs better? Why?
Note: More information about analysing queries you can find [here](#).

A:

[EDITABLE CODE]

```
--Version 1: Using category as a string ('documentary'):
EXPLAIN ANALYZE
SELECT m.title as "Movie Title", cat.name as "Category"
FROM movie m
INNER JOIN movie_Category mc on m.movie_id = mc.movie_id
INNER JOIN category cat ON mc.category_id = cat.category_id
WHERE cat.name = 'Documentary';

--Version 2: Using category ID ('Documentary' has category_id = 6):
EXPLAIN ANALYZE
```



```
SELECT m.title AS " Movie Title", cat.name AS "Category"
FROM movie m
INNER JOIN movie_Category mc ON m.movie_id = mc.movie_id
INNER JOIN category cat ON mc.category_id = cat.category_id
WHERE cat.category_id = 6;
```

[OUTPUT SCREENSHOT]

```
movie_rental=# select * from category;
category_id | name          | last_update
-----+-----+-----
1 | Action       | 2022-02-15 09:46:27+00
2 | Animation    | 2022-02-15 09:46:27+00
3 | Children     | 2022-02-15 09:46:27+00
4 | Classics     | 2022-02-15 09:46:27+00
5 | Comedy       | 2022-02-15 09:46:27+00
6 | Documentary   | 2022-02-15 09:46:27+00
7 | Drama        | 2022-02-15 09:46:27+00
8 | Family       | 2022-02-15 09:46:27+00
9 | Foreign      | 2022-02-15 09:46:27+00
10 | Games        | 2022-02-15 09:46:27+00
11 | Horror       | 2022-02-15 09:46:27+00
12 | Music        | 2022-02-15 09:46:27+00
13 | New          | 2022-02-15 09:46:27+00
14 | Sci-Fi       | 2022-02-15 09:46:27+00
15 | Sports       | 2022-02-15 09:46:27+00
16 | Travel       | 2022-02-15 09:46:27+00
(16 rows)
```

(category_id = 6 (Documentary))

```
movie_rental=# --Version 1: Using category as a string ('documentary'):
EXPLAIN ANALYZE
SELECT m.title as "Movie Title", cat.name as "Category"
FROM movie m
INNER JOIN movie_Category mc ON m.movie_id = mc.movie_id
INNER JOIN category cat ON mc.category_id = cat.category_id
WHERE cat.name = 'Documentary';
```

```
--Version 2: Using category ID (assuming 'Documentary' has category_id = 1):
EXPLAIN ANALYZE
SELECT m.title AS " Movie Title", cat.name AS "Category"
FROM movie m
INNER JOIN movie_Category mc ON m.movie_id = mc.movie_id
INNER JOIN category cat ON mc.category_id = cat.category_id
WHERE cat.category_id = 6;
```

QUERY PLAN

```
Nested Loop (cost=1.49..53.54 rows=62 width=133) (actual time=0.042..0.338 rows=68 loops=1)
-> Hash Join (cost=1.21..20.53 rows=62 width=122) (actual time=0.035..0.207 rows=68 loops=1)
    Hash Cond: (mc.category_id = cat.category_id)
    -> Seq Scan on movie_category mc (cost=0.00..16.00 rows=1000 width=8) (actual time=0.010..0.087 rows=1000 loops=1)
    -> Hash (cost=1.20..1.20 rows=1 width=122) (actual time=0.015..0.016 rows=1 loops=1)
        Buckets: 1024 Batches: 1 Memory Usage: 9kB
        -> Seq Scan on category cat (cost=0.00..1.20 rows=1 width=122) (actual time=0.009..0.010 rows=1 loops=1)
            Filter: ((name)::text = 'Documentary'::text)
            Rows Removed by Filter: 15
    -> Index Scan using movie_pkey on movie m (cost=0.28..0.53 rows=1 width=19) (actual time=0.002..0.002 rows=1 loops=68)
        Index Scan using (movie_id = mc.movie_id)
Planning Time: 0.327 ms
Execution Time: 0.382 ms
(13 rows)
```

QUERY PLAN

```
Nested Loop (cost=19.35..90.66 rows=68 width=133) (actual time=0.108..0.504 rows=68 loops=1)
-> Seq Scan on category cat (cost=0.00..1.20 rows=1 width=122) (actual time=0.007..0.008 rows=1 loops=1)
    Filter: (category_id = 6)
    Rows Removed by Filter: 15
-> Hash Join (cost=19.35..88.78 rows=68 width=19) (actual time=0.100..0.487 rows=68 loops=1)
    Hash Cond: (m.movie_id = mc.movie_id)
    -> Seq Scan on movie m (cost=0.00..65.00 rows=1000 width=19) (actual time=0.003..0.278 rows=1000 loops=1)
    -> Hash (cost=18.50..18.50 rows=68 width=8) (actual time=0.092..0.093 rows=68 loops=1)
        Buckets: 1024 Batches: 1 Memory Usage: 11kB
        -> Seq Scan on movie_category mc (cost=0.00..18.50 rows=68 width=8) (actual time=0.005..0.080 rows=68 loops=1)
            Filter: (category_id = 6)
            Rows Removed by Filter: 932
Planning Time: 0.184 ms
Execution Time: 0.525 ms
(14 rows)
```

Version 1 perform better and the reasons on why as follows:

1. **Indexing:** If an index is present on the name column in the category table, the database can quickly look up rows with 'Documentary' without scanning the whole table.
2. **Data Distribution:** If there are very few categories with the name 'Documentary' compared to categories with the ID '6', the first query could be faster.
3. **Query Planning and Optimization:** PostgreSQL's query planner may have chosen different strategies for executing the two queries based on table statistics, leading to different performance characteristics.

Q5. Using your SQL knowledge create a query that will list all movies that have original language Italian along with the release year. The movies should be listed by release year in chronological order.

A:

[EDITABLE CODE]

```
SELECT m.title, m.release_year
FROM movie m
INNER JOIN language l on l.language_id = m.language_id
WHERE m.original_language_id = 2
ORDER BY m.release_year;
```

[OUTPUT SCREENSHOT]

```
movie_rental=# SELECT m.title, m.release_year
FROM movie m
INNER JOIN language l on l.language_id = m.language_id
WHERE m.original_language_id = 2
ORDER BY m.release_year;
```

```

ARGONAUTS TOWN      2005
GASLIGHT CRUSADE    2005
DOZEN LION           2006
ILLUSION AMELIE      2006
PLUTO OLEANDER       2007
GLORY TRACY          2007
GOLDFINGER SENSIBILITY 2007
HOURS RAGE           2007
YOUNG LANGUAGE       2008
HANGING DEEP         2008
NECKLACE OUTBREAK    2009
FAMILY SWEET         2009
CLONES PINOCCHIO     2009
TOMORROW HUSTLER     2009
SIDE ARK             2009
CROOKED FROGMEN      2010
SPEAKEASY DATE       2010
POTTER CONNECTICUT   2010
WHISPERER GIANT      2010
HOLLOW JEOPARDY      2010
LOVER TRUMAN         2010
OZ LIAISONS          2011
EGG IGBY             2012
CITIZEN SHREK        2012
FERRIS MOTHER        2012
BRANNIGAN SUNRISE    2013
SLUMS DUCK           2014
PLATOON INSTINCT     2015
SHANE DARKNESS       2015
TRAFFIC HOBBIT       2015
CIDER DESIRE         2015
REMEMBER DIARY       2015
ELEMENT FREDDY       2015
OCTOBER SUBMARINE    2016
DESERT POSEIDON      2016
RAGING AIRPLANE      2016
MEMENTO ZOOLANDER    2017
SEA VIRGIN           2017
FLIGHT LIES          2018
ALTER VICTORY         2018
WILD APOLLO          2019
GRIT CLOCKWORK       2020
DOCTOR GRAIL         2021
HIGHBALL POTTER      2021
HARPER DYING         2021
KENTUCKIAN GIANT     2022
(153 rows)
(END)

```

Q6. List how many movies are in each category per original language. The output should be similar to the one below.

Category	Original Language	Number of Movies
Action	English	7
Action	French	17
Action	German	13
Action	Italian	5
Action	Japanese	6
Action	Mandarin	16
Animation	English	10
Animation	French	14
Animation	German	7
Animation	Italian	11

A:

[EDITABLE SCREENSHOT]

```
SELECT cat.name AS "Category", l2.name AS "Original Language", COUNT(*) AS
"Number of Movies"
FROM movie m
INNER JOIN language l ON m.language_id = l.language_id
INNER JOIN language l2 ON m.original_language_id = l2.language_id
INNER JOIN movie_Category mc ON m.movie_id = mc.movie_id
INNER JOIN category cat ON mc.category_id = cat.category_id
GROUP BY cat.name, l2.name
ORDER BY cat.name, l2.name;
```

[OUTPUT SCREENSHOT]

Category	Original Language	Number of Movies
Action	English	7
Action	French	17
Action	German	13
Action	Italian	5
Action	Japanese	6
Action	Mandarin	16
Animation	English	10
Animation	French	14
Animation	German	7
Animation	Italian	11
Animation	Japanese	14
Animation	Mandarin	10

Conclusion/Reflection

Through Lab 5, I deepened my understanding of SQL and its real-world applications. Working directly with the Movie Rental database, I executed and analyzed queries, which sharpened my comprehension of SQL syntax and its nuances. I was able to distinguish between the performance of different queries based on how they interact with the database's indexes, and the concept of sargability became clearer. In addition, by identifying and correcting SQL syntax errors, I gained more confidence in formulating my own SQL queries. Overall, Lab 5 provided valuable insights into how databases are managed and manipulated in real-world scenarios.

LAB 6 (21/07/2023) - Optimization

Lab Description

In Lab 6, advanced SQL abilities through the "movie_rental" database, focusing on data manipulation, query optimization, and index application for enhanced query performance. Key tasks included updating specific records, creating and dropping indexes, and creating a VIEW to simplify data retrieval.

- Q1. Connect to the movie_rental database (that you have installed in [LAB5](#)) and change the phone number from the current one to **02392844444** for address with id 100. Do not alter any other data/attribute.**

A:

[EDITABLE CODE]

```
-- Connect to 'movie_rental' database
\c movie_rental

-- Change address with id 100
UPDATE address
SET phone = '02392844444'
WHERE address_id = 100;

-- View updated address with id 100
SELECT * FROM address WHERE address_id = 100;
```

[OUTPUT SCREENSHOT]

```
up2200918=# \c movie_rental
You are now connected to database "movie_rental" as user "up2200918".
movie_rental=# UPDATE address
SET phone = '02392844444'
WHERE address_id = 100;
UPDATE 1
movie_rental=# SELECT * FROM address WHERE address_id = 100;
 address_id | address      | address2 | district | city_id | postal_code | phone      | last_update
-----+-----+-----+-----+-----+-----+-----+-----
      100 | 1308 Arecibo Way |         | Georgia  |      41 |      30695 | 02392844444 | 2023-07-21 10:49:31.694894+00
(1 row)
```

Q2. Create a query that will return all the details for the record where the phone number is **02392844444 and use EXPLAIN. Take a screenshot of the output.**

A:

[EDITABLE CODE]

```
EXPLAIN ANALYZE
SELECT *
FROM address a
WHERE phone = '02392844444';
```

[OUTPUT SCREENSHOT]

```
movie_rental=# SELECT *
FROM address a
WHERE phone = '02392844444';
 address_id | address          | address2 | district | city_id | postal_code | phone          | last_update
-----+-----+-----+-----+-----+-----+-----+-----
          100 | 1308 Arecibo Way |          | Georgia  |      41 |      30695 | 02392844444 | 2023-07-21 10:49:31.694894+00
(1 row)

movie_rental=# EXPLAIN ANALYZE
SELECT *
FROM address a
WHERE phone = '02392844444';
                                QUERY PLAN
-----
Seq Scan on address a  (cost=0.00..15.54 rows=1 width=63) (actual time=0.090..0.091 rows=1 loops=1)
  Filter: ((phone)::text = '02392844444'::text)
  Rows Removed by Filter: 602
  Planning Time: 0.079 ms
  Execution Time: 0.114 ms
(5 rows)
```

Q3. Apply an INDEX on the **phone attribute.**

A:

[EDITABLE CODE]

```
CREATE INDEX idx_address_phone ON address (phone);
```

[OUTPUT SCREENSHOT]

```
movie_rental=# CREATE INDEX idx_address_phone ON address (phone);
CREATE INDEX
```

Q4. Use EXPLAIN again. What are the differences? Which one performs better and why?

A:

- Without Index: In this case, the query may involve a "Sequential Scan" of the table, which means the database engine will go through all rows one by one to find the matching phone number. This could be less efficient, especially for large tables, as it requires reading the entire table.

- With Index: After applying the index, the query should utilize an "Index Scan." The index allows the database engine to quickly locate the rows with the specified phone number, resulting in a much faster execution time, particularly for tables with a large number of records.
- The **query with the index should perform significantly better** than the query without the index, especially when dealing with large datasets. The index provides an efficient way to narrow down the search and directly access the required rows, leading to improved performance in searching for the phone number '02392844444'.

Q5. Delete the INDEX you have created in Q3.

A:

[EDITABLE CODE]

```
DROP INDEX idx_address_phone;
```

[OUTPUT SCREENSHOT]

```
movie_rental=# DROP INDEX idx_address_phone;
DROP INDEX
```

Q6. Create a VIEW (*customer_details*) that will hold the customers details exactly as below.

Customer	Contact Details	Customer Address	Customer City	Customer Country
VERA MCCOY	VERA.MCCOY@sakilacustomer.org 886649065861	1168 Najafabad Parkway	Kabul	Afghanistan
MARIO CHEATHAM	MARIO.CHEATHAM@sakilacustomer.org 486784385440	1924 Shimonoseki Drive	Batna	Algeria
JUDY GRAY	JUDY.GRAY@sakilacustomer.org 107137400143	1031 Daugavpils Parkway	Bchar	Algeria
JUNE CARROLL	JUNE.CARROLL@sakilacustomer.org 506134035434	757 Rustenburg Avenue	Skikda	Algeria
ANTHONY SCHWAB	ANTHONY.SCHWAB@sakilacustomer.org 478229987054	1892 Naberezhnye Telny Lane	Tafuna	American Samoa
CLAUDE HERZOG	CLAUDE.HERZOG@sakilacustomer.org 105882218332	486 Ondo Parkway	Benguela	Angola
MARTIN BALES	MARTIN.BALES@sakilacustomer.org 106439158941	368 Hunuco Boulevard	Namibe	Angola
BOBBY BOUDREAU	BOBBY.BOUDREAU@sakilacustomer.org 934352415130	1368 Maracabo Boulevard	South Hill	Anguilla
JASON MORRISSEY	JASON.MORRISSEY@sakilacustomer.org 972574862516	1427 A Corua (La Corua) Place	Baha Blanca	Argentina
PERRY SWAFFORD	PERRY.SWAFFORD@sakilacustomer.org 914466027044	773 Dallas Manor	Quilmes	Argentina
DARRYL ASHCRAFT	DARRYL.ASHCRAFT@sakilacustomer.org 717566026669	166 Jinchang Street	Ezeiza	Argentina
MICHAEL FORMAN	MICHAEL.FORMAN@sakilacustomer.org 411549550611	203 Tambaram Street	Escobar	Argentina
JULIA FLORES	JULIA.FLORES@sakilacustomer.org 846225459260	1926 El Alto Avenue	La Plata	Argentina
ERIC ROBERT	ERIC.ROBERT@sakilacustomer.org 105470691550	430 Kumbakonam Drive	Santa F	Argentina
KIMBERLY LEE	KIMBERLY.LEE@sakilacustomer.org 934730187245	96 Tafuna Way	Crdoba	Argentina
LEONARD SCHOFIELD	LEONARD.SCHOFIELD@sakilacustomer.org 779461480495	88 Nagaon Manor	Tandil	Argentina
JORDAN ARCHULETA	JORDAN.ARCHULETA@sakilacustomer.org 817740355461	1229 Varanasi (Benares) Manor	Avellaneda	Argentina
LYDIA BURKE	LYDIA.BURKE@sakilacustomer.org 686015532180	1483 Pathankot Street	San Miguel de Tucuman	Argentina
WILLIE MARKHAM	WILLIE.MARKHAM@sakilacustomer.org 296394569728	1623 Kingstown Drive	Almirante Brown	Argentina
FLORENCE WOODS	FLORENCE.WOODS@sakilacustomer.org 330838016880	1532 Dzerzinsk Way	Merlo	Argentina
WILLIE HOMELL	WILLIE.HOMELL@sakilacustomer.org 991802825778	1244 Allappuzha (Alleppey) Place	Vicente Lopez	Argentina
STEPHANIE MITCHELL	STEPHANIE.MITCHELL@sakilacustomer.org 42384721397	42 Brindisi Place	Yerevan	Armenia
JILL HAWKINS	JILL.HAWKINS@sakilacustomer.org 931059836497	1440 Compton Place	Linz	Austria
AUDREY RAY	AUDREY.RAY@sakilacustomer.org 493008546074	1010 Klerksdorp Way	Graz	Austria
NORA HERRERA	NORA.HERRERA@sakilacustomer.org 621625204422	1587 Loja Manor	Salzburg	Austria

A:

[EDITABLE CODE]

```
CREATE VIEW customer_details AS
```

```

SELECT
    CONCAT(c.first_name, ' ', c.last_name) AS "Customer",
    CONCAT(c.email, ' | ', a.phone) AS "Contact Details",
    address AS "Customer Address",
    ci.city AS "Customer City",
    co.country AS "Customer Country"
FROM customer c
INNER JOIN address a ON c.address_id = a.address_id
INNER JOIN city ci ON a.city_id = ci.city_id
INNER JOIN country co ON ci.country_id = co.country_id
ORDER BY co.country;

```

[OUTPUT SCREENSHOT]

```

movie_rental=# CREATE VIEW customer_details AS
SELECT
    CONCAT(c.first_name, ' ', c.last_name) AS "Customer",
    CONCAT(c.email, ' | ', a.phone) AS "Contact Details",
    address AS "Customer Address",
    ci.city AS "Customer City",
    co.country AS "Customer Country"
FROM customer c
INNER JOIN address a ON c.address_id = a.address_id
INNER JOIN city ci ON a.city_id = ci.city_id
INNER JOIN country co ON ci.country_id = co.country_id
ORDER BY co.country;
CREATE VIEW
movie_rental=# select * from customer_details;

```

Customer	Contact Details	Customer Address	Customer City	Customer Country
VERA MCCOY	VERA.MCCOY@sakilacustomer.org 886649065861	1168 Najafabad Parkway	Kabul	Afghanistan
MARIO CHEATHAM	MARIO.CHEATHAM@sakilacustomer.org 406784385440	1924 Shimonoseki Drive	Batna	Algeria
JUDY GRAY	JUDY.GRAY@sakilacustomer.org 107137400143	1001 Daugavpils Parkway	Bchar	Algeria
JUNE CARROLL	JUNE.CARROLL@sakilacustomer.org 506134035434	757 Rustenburg Avenue	Skikda	Algeria
ANTHONY SCHWAB	ANTHONY.SCHWAB@sakilacustomer.org 478229987054	1892 Naberezhnye Teln Lane	Tafuna	American Samoa
CLAUDE HERZOG	CLAUDE.HERZOG@sakilacustomer.org 105882218332	486 Ondo Parkway	Benguela	Angola
MARTIN BALES	MARTIN.BALES@sakilacustomer.org 106439158941	368 Hunuco Boulevard	Namibe	Angola
BOBBY BOURDEAU	BOBBY.BOURDEAU@sakilacustomer.org 934352415130	1368 Manacabo Boulevard	South Hill	Anguilla
JASON MORRISSEY	JASON.MORRISSEY@sakilacustomer.org 972574862516	1427 A Corua (La Corua) Place	Baha Blanca	Argentina
PERRY SWAFFORD	PERRY.SWAFFORD@sakilacustomer.org 914466027044	773 Dallas Manor	Quilmes	Argentina
DARRYL ASHCRAFT	DARRYL.ASHCRAFT@sakilacustomer.org 717566026669	166 Jinchang Street	Ezeiza	Argentina
MICHAEL FORMAN	MICHAEL.FORMAN@sakilacustomer.org 411549550611	203 Tambaram Street	Escobar	Argentina
JULIA FLORES	JULIA.FLORES@sakilacustomer.org 846225459260	1926 El Alto Avenue	La Plata	Argentina
ERIC ROBERT	ERIC.ROBERT@sakilacustomer.org 105470691550	430 Kumbakonam Drive	Santa F	Argentina
KIMBERLY LEE	KIMBERLY.LEE@sakilacustomer.org 934730187245	96 Tafuna Way	Crdoba	Argentina
LEONARD SCHOFFELD	LEONARD.SCHOFFELD@sakilacustomer.org 778461480495	88 Nagoon Manor	Tandil	Argentina
JORDAN ARCHULETA	JORDAN.ARCHULETA@sakilacustomer.org 817740355461	1229 Varanasi (Benares) Manor	Avellaneda	Argentina
LYDIA BURKE	LYDIA.BURKE@sakilacustomer.org 686015532180	1483 Pathankot Street	San Miguel de Tucum	Argentina
WILLIE MARKHAM	WILLIE.MARKHAM@sakilacustomer.org 296394569728	1623 Kingstown Drive	Almirante Brown	Argentina
FLORENCE WOODS	FLORENCE.WOODS@sakilacustomer.org 330838016880	1532 Dzerzinsk Way	Merlo	Argentina
WILLIE HOWELL	WILLIE.HOWELL@sakilacustomer.org 991802825778	1244 Allappuzha (Alleppey) Place	Vicente Lpez	Argentina
STEPHANIE MITCHELL	STEPHANIE.MITCHELL@sakilacustomer.org 42384721397	42 Brindisi Place	Yerevan	Armenia
JILL HAWKINS	JILL.HAWKINS@sakilacustomer.org 931059836497	1440 Compton Place	Linz	Austria
AUDREY RAY	AUDREY.RAY@sakilacustomer.org 493008546074	1010 Klerksdorp Way	Graz	Austria
NORA HERRERA	NORA.HERRERA@sakilacustomer.org 621625204422	1587 Loja Manor	Salzburg	Austria
ANDREW PURDY	ANDREW.PURDY@sakilacustomer.org 119501405123	431 Szekesfehrv Avenue	Baku	Azerbaijan
RAYMOND MCWHORTER	RAYMOND.MCWHORTER@sakilacustomer.org 834626715837	503 Sogamoso Loop	Sumgayit	Azerbaijan
SETH HANNON	SETH.HANNON@sakilacustomer.org 864392582257	1759 Niznekamsk Avenue	al-Manama	Bahrain
STEPHEN QUALLS	STEPHEN.QUALLS@sakilacustomer.org 38988715447	1838 Tabriz Lane	Dhaka	Bangladesh
MICHELLE CLARK	MICHELLE.CLARK@sakilacustomer.org 892775750063	262 A Corua (La Corua) Parkway	Tangail	Bangladesh
FRANK WAGGONER	FRANK.WAGGONER@sakilacustomer.org 965273813662	1816 Bydgoszcz Loop	Jamalnur	Bangladesh
CLARA SHAW	CLARA.SHAW@sakilacustomer.org 563660187896	1027 Songkhla Manor	Mlodetno	Belarus
CORY MEEHAN	CORY.MEEHAN@sakilacustomer.org 338244023543	556 Asuncin Way	Mogiljov	Belarus
JOEL FRANCISCO	JOEL.FRANCISCO@sakilacustomer.org 82619513349	287 Cuautla Boulevard	Sucre	Bolivia
JON WILES	JON.WILES@sakilacustomer.org 205524798287	659 Gatineau Boulevard	El Alto	Bolivia
ELLEN SIMPSON	ELLEN.SIMPSON@sakilacustomer.org 582835362905	1666 Qomshah Drive	Po	Brazil
KENT ARSENAULT	KENT.ARSENAULT@sakilacustomer.org 410877354933	32 Liaocheng Way	Juiz de Fora	Brazil
TINA SIMMONS	TINA.SIMMONS@sakilacustomer.org 132986892228	984 Effen-Alaiye Avenue	Goinia	Brazil
CHARLOTTE HUNTER	CHARLOTTE.HUNTER@sakilacustomer.org 935448624185	758 Junan Lane	guas Lndas de Gois	Brazil
JESUS MCCARTNEY	JESUS.MCCARTNEY@sakilacustomer.org 419009857119	259 Ipho Drive	Guaruj	Brazil
ANTONIO MEEK	ANTONIO.MEEK@sakilacustomer.org 841876514789	1190 0 Place	Bag	Brazil
FRANCISCO SKIDMORE	FRANCISCO.SKIDMORE@sakilacustomer.org 876491807547	614 Denizli Parkway	So Leopoldo	Brazil
CLAYTON BARBEE	CLAYTON.BARBEE@sakilacustomer.org 380077794770	1407 Pachuca de Soto Place	Alvorada	Brazil
NATALIE MEYER	NATALIE.MEYER@sakilacustomer.org 873492228462	1201 Qomshah Manor	Aparecida de Goinia	Brazil
JOSEPH JOY	JOSEPH.JOY@sakilacustomer.org 573441301529	1354 Siegen Street	Angra dos Reis	Brazil
TIMOTHY BUNN	TIMOTHY.BUNN@sakilacustomer.org 829116184079	981 Kumbakonam Place	Braslia	Brazil

Q7. Using a subquery list all cities from the United Kingdom along with the city ID.

A:

[EDITABLE CODE]

```
SELECT city_id, city
FROM city
WHERE country_id = (
    SELECT country_id
    FROM country
    WHERE country = 'United Kingdom'
);
```

[OUTPUT SCREENSHOT]

```
movie_rental=# SELECT city_id, city
FROM city
WHERE country_id = (
    SELECT country_id
    FROM country
    WHERE country = 'United Kingdom'
);
city_id |      city
-----+-----
      88 | Bradford
     149 | Dundee
     312 | London
     494 | Southampton
     495 | Southend-on-Sea
     496 | Southport
     500 | Stockport
     589 | York
(8 rows)
```

Q8. As an extension of the previous query, list the cities from the United Kingdom and France using the **IN** operator, this time showing the country name as well. Group them by country.

A:

[EDITABLE CODE]

```
SELECT c.city_id, c.city, co.country AS country_name
FROM city c
INNER JOIN country co ON c.country_id = co.country_id
WHERE co.country IN ('United Kingdom', 'France')
ORDER BY co.country;
```

[OUTPUT SCREENSHOT]

```

movie_rental=# SELECT c.city_id, c.city, co.country AS country_name
FROM city c
INNER JOIN country co ON c.country_id = co.country_id
WHERE co.country IN ('United Kingdom', 'France')
ORDER BY co.country;
 city_id |      city      | country_name
-----+-----+-----
      92 | Brest          | France
     543 | Toulon         | France
     544 | Toulouse       | France
     297 | Le Mans        | France
     495 | Southend-on-Sea | United Kingdom
     496 | Southport      | United Kingdom
     500 | Stockport      | United Kingdom
      88 | Bradford       | United Kingdom
     589 | York           | United Kingdom
     149 | Dundee         | United Kingdom
     312 | London         | United Kingdom
     494 | Southampton    | United Kingdom
(12 rows)

```

Conclusion/Reflection

The practical application of SQL concepts deepened my understanding, particularly in the areas of data manipulation, query optimization, and index usage. I gained valuable insights into the effective use of indexes for performance improvement, especially in large datasets. Moreover, the creation and use of a VIEW simplified data access, reinforcing the utility of these structures in SQL. Additionally, the use of subqueries and the IN operator to filter data from multiple countries further expanded my SQL capabilities.

LAB 8 (22/07/2023) -Data Manipulation and Query Output Formatting

Lab Description

In Lab 8, expanded proficiency in SQL through the "movie_rental" database, focusing on complex querying across multiple tables and data manipulation. The lab tasks involved pattern matching using ILIKE/LIKE, wildcard characters, counting unique values, and sorting data.

Additionally, addressed scenarios that involved handling data across multiple rows, such as identifying the most frequent occurrence of a value.

For this week we will use the movie_rental database. Make sure that you have the [database installed](#) and you are connected to it. You can use [this ERD](#) and this interactive [Data Dictionary](#).

Q1. List **id, first name and last name** of all actors that have the first name as **Scarlett**. The query should ignore the name capitalisation. Hint: *Look for the **ILIKE/LIKE** keywords or wildcard characters.*

A:

[EDITABLE CODE]

```
SELECT actor_id, first_name, last_name
FROM actor
WHERE first_name ILIKE 'Scarlett';
```

[OUTPUT SCREENSHOT]

```
movie_rental=# SELECT actor_id, first_name, last_name
FROM actor
WHERE first_name ILIKE 'Scarlett';
 actor_id | first_name | last_name 
-----+-----+-----
       81 | SCARLETT  | DAMON
      124 | SCARLETT  | BENING
(2 rows)
```

Q2. How many *unique last names* are in actors' names? (e.g. If they are 3 Smith's will be counted only once).

A:

[EDITABLE CODE]

```
SELECT COUNT(DISTINCT last_name) AS unique_last_names
FROM actor;
```

[OUTPUT SCREENSHOT]

```
movie_rental=# SELECT COUNT(DISTINCT last_name) AS unique_last_names
FROM actor;
 unique_last_names
-----
                121
(1 row)
```

Q3. Following from the previous query, we know how many unique last names are in the database. However, how many are truly unique?

List in alphabetical order the last names that are repeated only once (e.g. if a last name is more than once in the database it will not be considered).

A:

[EDITABLE CODE]

```
SELECT last_name
FROM actor
GROUP BY last_name
HAVING COUNT(last_name) = 1
ORDER BY last_name;
```

[OUTPUT SCREENSHOT]

```

movie_rental=# SELECT last_name
FROM actor
GROUP BY last_name
HAVING COUNT(last_name) = 1
ORDER BY last_name;

```

```

last_name
-----
ASTAIRE
BACALL
BALE
BALL
BARRYMORE
BASINGER
BERGEN
BERGMAN
BIRCH
BLOOM
BRIDGES
BULLOCK
CARREY
CHAPLIN
CLOSE
COSTNER
CROWE
CRUISE
CRUZ
DAMON
DAY-LEWIS
DERN
DREYFUSS
DUNST
GABLE
GIBSON
GOLDBERG
GRANT
HAWKE
HESTON
HOPE
HUDSON
HUNT
HURT
JOLIE
JOVOVICH
LEIGH
LOLLOBRIGIDA
MALDEN
MANSFIELD
MARX
MCDORMAND
MIRANDA
NICHOLSON
PESCI
PFEIFFER
:

```

```

DAY-LEWIS
DERN
DREYFUSS
DUNST
GABLE
GIBSON
GOLDBERG
GRANT
HAWKE
HESTON
HOPE
HUDSON
HUNT
HURT
JOLIE
JOVOVICH
LEIGH
LOLLOBRIGIDA
MALDEN
MANSFIELD
MARX
MCDORMAND
MIRANDA
NICHOLSON
PESCI
PFEIFFER
PHOENIX
PINKETT
PITT
POSEY
PRESLEY
REYNOLDS
RYDER
SINATRA
SOBIESKI
STALLONE
SUVARI
SWANK
TAUTOU
TOMEI
VOIGHT
WALKEN
WAYNE
WILSON
WITHERSPOON
WRAY
(66 rows)

(END)

```

Q4. List the **first name and last name** of the actor that appear in most movies and the **number of movies**. Hint: Look for **USING**.

A:

[EDITABLE CODE]

```
SELECT a.first_name, a.last_name, COUNT(ma.movie_id) AS movie_count
FROM actor a
JOIN movie_actor ma USING (actor_id)
GROUP BY a.actor_id, a.first_name, a.last_name
ORDER BY movie_count DESC
LIMIT 1;
```

[OUTPUT SCREENSHOT]

```
movie_rental=# SELECT a.first_name, a.last_name, COUNT(ma.movie_id) AS movie_count
FROM actor a
JOIN movie_actor ma USING (actor_id)
GROUP BY a.actor_id, a.first_name, a.last_name
ORDER BY movie_count DESC
LIMIT 1;
 first_name | last_name | movie_count
-----+-----+-----
GINA       | DEGENERES |          42
(1 row)
```

Q5. Each store has several copies of each movie. List the available **store 1 inventory** for the movie named '**Purple Movie**'. How many copies are available and what are the inventory IDs?

A:

[EDITABLE CODE]

```
SELECT i.inventory_id, COUNT(*) AS copies_available
FROM inventory i
JOIN movie m ON i.movie_id = m.movie_id
JOIN store s ON i.store_id = s.store_id
WHERE m.title = 'PURPLE MOVIE' AND i.store_id = 1
GROUP BY i.inventory_id;
```

[OUTPUT SCREENSHOT]

```

movie_rental=# SELECT i.inventory_id, COUNT(*) AS copies_available
FROM inventory i
JOIN movie m ON i.movie_id = m.movie_id
JOIN store s ON i.store_id = s.store_id
WHERE m.title = 'PURPLE MOVIE' AND i.store_id = 1
GROUP BY i.inventory_id;
 inventory_id | copies_available
-----+-----
          3207 |                1
          3208 |                1
          3209 |                1
          3210 |                1
(4 rows)

```

- Q6.** List staff name and last name, along with their home address, city and email address. The data should be presented in a nice format (exactly as below).

Staff Name	Staff Address	Staff City	Staff Email
Mike Hillyer	23 Workhaven Lane	Lethbridge	mike.hillyer@sakilastaff.com
Jon Stephens	1411 Lillydale Drive	Woodridge	jon.stephens@sakilastaff.com

A:

[EDITABLE CODE]

```

SELECT CONCAT(first_name, ' ', last_name) AS "Staff Name",
       address AS "Staff Address",
       city AS "Staff City",
       email AS "Staff Email"
FROM staff s
INNER JOIN address a on s.address_id = a.address_id
INNER JOIN city ci on a.city_id = ci.city_id;

```

[OUTPUT SCREENSHOT]

```

up2200918=# \c movie_rental
You are now connected to database "movie_rental" as user "up2200918".
movie_rental=# SELECT CONCAT(first_name, ' ', last_name) AS "Staff Name",
       address AS "Staff Address",
       city AS "Staff City",
       email AS "Staff Email"
FROM staff s
INNER JOIN address a on s.address_id = a.address_id
INNER JOIN city ci on a.city_id = ci.city_id;
 Staff Name | Staff Address | Staff City | Staff Email
-----+-----+-----+-----
Mike Hillyer | 23 Workhaven Lane | Lethbridge | Mike.Hillyer@sakilastaff.com
Jon Stephens | 1411 Lillydale Drive | Woodridge | Jon.Stephens@sakilastaff.com
(2 rows)

```

Q7. List the *name and the last name* in alphabetical order (on surname) of all actors that have acted in the movie named '**Agent Truman**'. The names should appear in a single column just the name of the actors, not the movie.

A:

[EDITABLE CODE]

```
SELECT CONCAT(a.first_name, ' ', a.last_name) AS "Actor Name"
FROM actor a
JOIN movie_actor ma on a.actor_id = ma.actor_id
JOIN movie m on ma.movie_id = m.movie_id
WHERE m.title = 'AGENT TRUMAN'
ORDER BY last_name, first_name;
```

[OUTPUT SCREENSHOT]

```
movie_rental=# SELECT CONCAT(a.first_name, ' ', a.last_name) AS "Actor Name"
FROM actor a
JOIN movie_actor ma on a.actor_id = ma.actor_id
JOIN movie m on ma.movie_id = m.movie_id
WHERE m.title = 'AGENT TRUMAN'
ORDER BY last_name, first_name;
   Actor Name
-----
KENNETH HOFFMAN
SANDRA KILMER
JAYNE NEESON
WARREN NOLTE
KIRSTEN PALTROW
REESE WEST
MORGAN WILLIAMS
(7 rows)
```

Conclusion/Reflection

Lab 8 was an engaging learning experience that allowed us to apply and enhance our SQL skills in a practical context. I got to explore advanced SQL concepts, such as pattern matching, counting unique instances, and ordering data. I also gained valuable insights into how to tackle complex queries, handle data distributed across multiple rows, and present data in a user-friendly format.

This lab further solidified my SQL abilities, enabling me to analyze complex data patterns and optimize database querying. This knowledge will be beneficial for handling database-related tasks and projects in the future.

LAB 9 (25/07/2023) - Advanced SQL

Lab Description

In Lab 9, we continued to explore the "movie_rental" database with a specific focus on database functions and stored procedures. We learned how to list all the functions available in the database, call a function with parameters to retrieve specific data, and create and utilize a stored procedure for data insertion. The lab also involved tasks such as modifying a column data type, ensuring uniqueness, and performing calculations to derive insights from the data.

For this week we will use the movie_rental database. Make sure that you have the [database installed](#) and you are connected to it. You can use [this ERD](#) and [Data Dictionary](#).

Q1. List all the functions available in the movie_rental database.

A:

[EDITABLE CODE]

```
SELECT routine_name
FROM information_schema.routines
WHERE specific_schema = 'public'
AND routine_type = 'FUNCTION';
```

[OUTPUT SCREENSHOT]

```
movie_rental=# SELECT routine_name
FROM information_schema.routines
WHERE specific_schema = 'public'
AND routine_type = 'FUNCTION';
 routine_name
-----
group_concat
movie_in_stock
movie_not_in_stock
last_day
inventory_in_stock
get_customer_balance
inventory_held_by_customer
last_updated
rewards_report
(9 rows)
```

Q2. In the database is a function (*movie_in_stock*), with two INT parameters (*p_movie_id* and *p_store_id*), that will return all available copies of a given movie in a particular store. Using the named function, find the number of copies of the movie Named “*Angels Life*” in store 1.

Hint: Look for [named notation](#) syntax format to call a function parameter.

The output should look like below.

```
+-----+
| Movies in Stock |
+-----+
| 124             |
| 125             |
| 126             |
| 127             |
+-----+
SELECT 4
```

A:

[EDITABLE CODE]

```
-- find movie_id where title is Angels Life
SELECT movie_id FROM movie WHERE title = 'ANGELS LIFE';
-- output movie_id 25
SELECT movie_in_stock(p_movie_id := 25, p_store_id := 1) AS "Movies In
Stock";
```

[OUTPUT SCREENSHOT]

```
movie_rental=# -- find movie_id where title is Angels Life
SELECT movie_id FROM movie WHERE title = 'ANGELS LIFE';
-- output movie_id 25
SELECT movie_in_stock(p_movie_id := 25, p_store_id := 1) AS "Movies In Stock";
 movie_id
-----
        25
(1 row)

Movies In Stock
-----
        124
        125
        126
        127
(4 rows)
```

Q3. Create a [stored procedure](#) (sp_add_new_actor) that will automatically insert a new actor. Insert your own name through stored procedure **CALL** e.g. **CALL sp_add_new_actor('Val', 'Adamescu');**

You don't need to allocate any ID as is automatically added by SERIAL PK and the last_update is automatically updated by another function. The procedure should take only two parameters (First Name & Last Name).

A:

[EDITABLE CODE]

```
-- Stored procedure
CREATE OR REPLACE PROCEDURE sp_add_new_actor(IN first_name TEXT, IN
last_name TEXT)
LANGUAGE plpgsql
AS $$
BEGIN
    INSERT INTO actor(first_name, last_name)
    VALUES (first_name, last_name);
END;
$$;

-- Call stored procedure
CALL sp_add_new_actor('KJ', 'NG');

-- Check whether if KJ NG actor is inserted in actor table
SELECT *
FROM actor
WHERE first_name = 'KJ' AND last_name = 'NG';
```

[OUTPUT SCREENSHOT]

```
movie_rental=# -- Stored procedure
CREATE OR REPLACE PROCEDURE sp_add_new_actor(IN first_name TEXT, IN last_name TEXT)
LANGUAGE plpgsql
AS $$
BEGIN
    INSERT INTO actor(first_name, last_name)
    VALUES (first_name, last_name);
END;
$$;

-- Call stored procedure
CALL sp_add_new_actor('KJ', 'NG');

-- Check whether if KJ NG actor is inserted in actor table
SELECT *
FROM actor
WHERE first_name = 'KJ' AND last_name = 'NG';
CREATE PROCEDURE
CALL
 actor_id | first_name | last_name | last_update
-----+-----+-----+-----
      206 | KJ        | NG        | 2023-07-22 06:58:36.401896+00
(1 row)

movie_rental=#
```

- Q4.** Create a new column in the country table named country_code as VARCHAR.

A:

[EDITABLE CODE]

```
ALTER TABLE country
ADD COLUMN country_code VARCHAR;
```

[OUTPUT SCREENSHOT]

```
movie_rental=# ALTER TABLE country
ADD COLUMN country_code VARCHAR;
ALTER TABLE
```

- Q5.** oops, my bad.. CHAR(2) data type will be more efficient. Modify the country_code from VARCHAR to CHAR(2) and to have only unique values for the column country_code. Use \d country and take a screenshot of the table details.

A:

[EDITABLE CODE]

```
-- Step 1: Change the data type of the country_code column to CHAR(2)
ALTER TABLE country
ALTER COLUMN country_code TYPE CHAR(2);

-- Step 2: Update any existing values to ensure they are exactly 2
characters long
UPDATE country
SET country_code = SUBSTRING(country_code FROM 1 FOR 2);

-- Step 3: Add a unique constraint on the country_code column
ALTER TABLE country
ADD CONSTRAINT country_code_unique UNIQUE (country_code);
```

[OUTPUT SCREENSHOT]

```
movie_rental=# ALTER TABLE country
ADD COLUMN country_code VARCHAR;
ALTER TABLE
movie_rental=# -- Step 1: Change the data type of the country_code column to CHAR(2)
ALTER TABLE country
ALTER COLUMN country_code TYPE CHAR(2);

-- Step 2: Update any existing values to ensure they are exactly 2 characters long
UPDATE country
SET country_code = SUBSTRING(country_code FROM 1 FOR 2);

-- Step 3: Add a unique constraint on the country_code column
ALTER TABLE country
ADD CONSTRAINT country_code_unique UNIQUE (country_code);
ALTER TABLE
UPDATE 109
ALTER TABLE
movie_rental=# \d country
```

Column	Type	Collation	Nullable	Default
country_id	integer		not null	nextval('country_country_id_seq'::regclass)
country	character varying(100)		not null	
last_update	timestamp with time zone		not null	now()
country_code	character(2)			

```
Indexes:
    "country_pkey" PRIMARY KEY, btree (country_id)
    "country_code_unique" UNIQUE CONSTRAINT, btree (country_code)
Referenced by:
    TABLE "city" CONSTRAINT "city_country_id_fkey" FOREIGN KEY (country_id) REFERENCES country(country_id) ON UPDATE CASCADE ON DELETE RESTRICT
Triggers:
    last_updated BEFORE UPDATE ON country FOR EACH ROW EXECUTE FUNCTION last_updated()
```

Q6. Now if everything is ready, insert the country code as **UK** for the United Kingdom and create the output as country id, country name and country code.

A:

[EDITABLE CODE]

```
-- Insert country_name, and country_code.
INSERT INTO country (country, country_code)
VALUES ('United Kingdom', 'UK');

-- Select country id, name and code
SELECT country_id, country, country_code
FROM country;
```

[OUTPUT SCREENSHOT]

```
movie_rental=# -- Insert country_name, and country_code.
INSERT INTO country (country, country_code)
VALUES ('United Kingdom', 'UK');
-- Select country id, name and code
SELECT country_id, country, country_code
FROM country;
INSERT 0 1
```

country_id	country	country_code
110	United Kingdom	UK
1	Afghanistan	
2	Algeria	
3	American Samoa	
4	Angola	
5	Anguilla	
6	Argentina	
7	Armenia	
8	Australia	
9	Austria	
10	Azerbaijan	
11	Bahrain	
12	Bangladesh	
13	Belarus	
14	Bolivia	
15	Brazil	
16	Brunei	
17	Bulgaria	
18	Cambodia	
19	Cameroon	
20	Canada	
21	Chad	
22	Chile	
23	China	
24	Colombia	
25	Congo, The Democratic Republic of the	
26	Czech Republic	
27	Dominican Republic	
28	Ecuador	
29	Egypt	
30	Estonia	
31	Ethiopia	
32	Faroe Islands	
33	Finland	
34	France	
35	French Guiana	
36	French Polynesia	
37	Gambia	
38	Germany	
39	Greece	
40	Greenland	
41	Holy See (Vatican City State)	
42	Hong Kong	
43	Hungary	
44	India	
45	Indonesia	

Q7. We are planning some migration of our data but we don't want to transfer everything. Create a new table (new_staff) and copy only id, first & last name and email address.

A:

[EDITABLE CODE]

```
CREATE TABLE new_staff (
    new_staff_id SERIAL PRIMARY KEY,
    first_name TEXT,
    last_name TEXT,
    email TEXT
);
```

[OUTPUT SCREENSHOT]

```
movie_rental=# CREATE TABLE new_staff (
    new_staff_id SERIAL PRIMARY KEY,
    first_name TEXT,
    last_name TEXT,
    email TEXT
);
CREATE TABLE
```

Q8. We have so many copies of each movie for rental but how many? The output of the movies in inventory must be exactly as below. Column naming, text formatting, number of copies starting from the movies with most copies until the ones with less copies.

Movie Title	Number of Copies
ACADEMY DINOSAUR	8
APACHE DIVINE	8
BEVERLY OUTLAW	8
BINGO TALENTED	8
BOOGIE AMELIE	8
BOUND CHEAPER	8
BUCKET BROTHERHOOD	8
BUTTERFLY CHOCOLAT	8
CAT CONEHEADS	8
CONFIDENTIAL INTERVIEW	8
CROSSROADS CASUALTIES	8
CURBOARD SINNERS	8
CURTAIN VIDEOTAPE	8
DANCING FEVER	8
DEER VIRGINIAN	8
DINOSAUR SECRETARY	8
DOGMA FAMILY	8
DYNAMITE TARZAN	8
EXPENDABLE STALLION	8
FAMILY SWEET	8
FORWARD TEMPLE	8
FROST HEAD	8
GARDEN ISLAND	8
GIANT TROOPERS	8

SOLDIERS EVOLUTION	2
SOUP WISDOM	2
SPEED SUIT	2
STALLION SUNDANCE	2
SUNSET RACER	2
TARZAN VIDEOTAPE	2
TEQUILA PAST	2
TEXAS WATCH	2
TRAFFIC HOBBIT	2
TRAIN BUNCH	2
TREATMENT JEKYLL	2
UNTOUCHABLES SUNRISE	2
VANISHED GARDEN	2
VISION TORQUE	2
WARLOCK WEREWOLF	2
WATERSHIP FRONTIER	2
WILD APOLLO	2
WONDERFUL DROP	2
WORLD LEATHERNECKS	2
YOUNG LANGUAGE	2
YOUTH KICK	2
ZHIVAGO CORE	2

SELECT 958	
(END)	

A:

[EDITABLE CODE]

```
SELECT m.title AS "Movie Title",
       COUNT(*) AS "Number of Copies"
FROM movie m
JOIN inventory i ON m.movie_id = i.movie_id
GROUP BY m.movie_id, m.title
ORDER BY COUNT(*) DESC, "Movie Title" ASC;
```

[OUTPUT SCREENSHOT]

```
movie_rental=# SELECT m.title AS "Movie Title", COUNT(*) AS "Number of Copies"
FROM movie m
JOIN inventory i ON m.movie_id = i.movie_id
GROUP BY m.movie_id, m.title
ORDER BY COUNT(*) DESC, "Movie Title" ASC;
movie_rental=#
```

Movie Title	Number of Copies
ACADEMY DINOSAUR	8
APACHE DIVINE	8
BEVERLY OUTLAW	8
BINGO TALENTED	8
BOOGIE AMELIE	8
BOUND CHEAPER	8
BUCKET BROTHERHOOD	8
BUTTERFLY CHOCOLAT	8
CAT CONEHEADS	8
CONFIDENTIAL INTERVIEW	8
CROSSROADS CASUALTIES	8
CUPBOARD SINNERS	8
CURTAIN VIDEOTAPE	8
DANCING FEVER	8
DEER VIRGINIAN	8
DINOSAUR SECRETARY	8
DOGMA FAMILY	8
DYNAMITE TARZAN	8
EXPENDABLE STALLION	8
FAMILY SWEET	8
FORWARD TEMPLE	8
FROST HEAD	8
GARDEN ISLAND	8
GIANT TROOPERS	8
GILMORE BOILED	8
GLEAMING JAMBREAKER	8
GOODFELLAS SALUTE	8
GREATEST NORTH	8
GRIT CLOCKWORK	8
HARRY IDAHO	8
HEAVYWEIGHTS BEAST	8
HOBBIT ALIEN	8
HORROR REIGN	8
HUSTLER PARTY	8
INNOCENT USUAL	8
INVASION CYCLONE	8
JUGGLER HARDLY	8
KISS GLORY	8
LOATHING LEGALLY	8
LOSE INCH	8
MARRIED GO	8
METROPOLIS COMA	8
MOCKINGBIRD HOLLYWOOD	8
MOON BUNCH	8
MUSCLE BRIGHT	8
NETWORK PEAK	8

MULHOLLAND BEAST	2
MUSSOLINI SPOILERS	2
MYSTIC TRUMAN	2
NEWSIES STORY	2
OKLAHOMA JUMANJI	2
PANIC CLUB	2
PAPI NECKLACE	2
PARK CITIZEN	2
PHANTOM GLORY	2
PIZZA JUMANJI	2
PLATOON INSTINCT	2
PRESIDENT BANG	2
PRIVATE DROP	2
PUNK DIVORCE	2
REBEL AIRPORT	2
RECORDS ZORRO	2
REDS POCUS	2
RUNAWAY TENENBAUMS	2
RUSHMORE MERMAID	2
SCHOOL JACKET	2
SENSE GREEK	2
SEVEN SWARM	2
SIMON NORTH	2
SLING LUKE	2
SOLDIERS EVOLUTION	2
SOUP WISDOM	2
SPEED SUIT	2
STALLION SUNDANCE	2
SUNSET RACER	2
TARZAN VIDEOTAPE	2
TEQUILA PAST	2
TEXAS WATCH	2
TRAFFIC HOBBIT	2
TRAIN BUNCH	2
TREATMENT JEKYLL	2
UNTOUCHABLES SUNRISE	2
VANISHED GARDEN	2
VISION TORQUE	2
WARLOCK WEREWOLF	2
WATERSHIP FRONTIER	2
WILD APOLLO	2
WONDERFUL DROP	2
WORLD LEATHERNECKS	2
YOUNG LANGUAGE	2
YOUTH KICK	2
ZHIVAGO CORE	2
(958 rows)	

(END)

Q9. What is the average movie length per category? Round it to the nearest two decimal places in descending order as below.

Category	Average movie length in Minutes
Sports	128.20
Games	127.84
Foreign	121.70
Drama	120.84
Comedy	115.83
Family	114.78
Music	113.65
Travel	113.32
Horror	112.48
Classics	111.67
Action	111.61
New	111.13
Animation	111.02
Children	109.80
Documentary	108.75
Sci-Fi	108.20

A:

[EDITABLE CODE]

```
SELECT cat.name AS "Category",
       ROUND(AVG(m.length), 2) AS "Average movie length in Minutes"
FROM movie m
JOIN movie_category mc ON m.movie_id = mc.movie_id
JOIN category cat ON mc.category_id = cat.category_id
GROUP BY cat.name
ORDER BY "Average movie length in Minutes" DESC;
```

[OUTPUT SCREENSHOT]

```
movie_rental=# SELECT cat.name AS "Category",
       ROUND(AVG(m.length), 2) AS "Average movie length in Minutes"
FROM movie m
JOIN movie_category mc ON m.movie_id = mc.movie_id
JOIN category cat ON mc.category_id = cat.category_id
GROUP BY cat.name
ORDER BY "Average movie length in Minutes" DESC;
 Category | Average movie length in Minutes
-----+-----
 Sports   | 128.20
 Games    | 127.84
 Foreign  | 121.70
 Drama    | 120.84
 Comedy   | 115.83
 Family   | 114.78
 Music    | 113.65
 Travel   | 113.32
 Horror   | 112.48
 Classics | 111.67
 Action   | 111.61
 New      | 111.13
 Animation| 111.02
 Children | 109.80
 Documentary | 108.75
 Sci-Fi   | 108.20
(16 rows)
```


Q10. We know that the average length of all movies is 115.27. Which categories have movies above average? Do not use LIMIT but select only categories that are above the average in descending order and rounded to nearest two decimal places.
Hint: This query can be achieved in multiple ways. You can use HAVING, USING, and subqueries.

A:

[EDITABLE CODE]

```
SELECT cat.name AS "Category",
       ROUND(AVG(m.length), 2) AS "Average movie length in Minutes"
FROM movie m
JOIN movie_category mc ON m.movie_id = mc.movie_id
JOIN category cat ON mc.category_id = cat.category_id
GROUP BY cat.name
HAVING ROUND(AVG(m.length), 2) > 115.27
ORDER BY ROUND(AVG(m.length), 2) DESC;
```

[OUTPUT SCREENSHOT]

```
movie_rental=# SELECT cat.name AS "Category",
       ROUND(AVG(m.length), 2) AS "Average movie length in Minutes"
FROM movie m
JOIN movie_category mc ON m.movie_id = mc.movie_id
JOIN category cat ON mc.category_id = cat.category_id
GROUP BY cat.name
HAVING ROUND(AVG(m.length), 2) > 115.27
ORDER BY ROUND(AVG(m.length), 2) DESC;
Category | Average movie length in Minutes
-----+-----
Sports   | 128.20
Games    | 127.84
Foreign  | 121.70
Drama    | 120.84
Comedy   | 115.83
(5 rows)
```

Conclusion/Reflection

Lab 9 provided an excellent opportunity to dive deeper into the practical aspects of SQL, particularly functions and stored procedures. I learned how to use existing functions to retrieve specific data and how to create a stored procedure to automate common database operations, such as inserting a new actor.

The lab improved my understanding of how to effectively use these powerful SQL features to simplify database tasks and improve code reusability. Moreover, working on modifying column data types and ensuring uniqueness reinforced my understanding of database design principles and the importance of data integrity.

Overall, Lab 9 was a challenging but enlightening experience, significantly enhancing my SQL skills and practical knowledge of databases. This knowledge will be invaluable in future work involving database manipulation and data analysis.

LAB 10 (25/07/2023) - Monitoring, Tuning, Backup & Recovery

Lab Description

Lab 10 was an intensive hands-on exercise that dove into the nitty-gritty of database administration through Linux-specific commands on the "movie_rental" database. The lab guided us through the creation of a backup for the database, followed by its deletion and subsequent restoration from the backup. Alongside this, we also executed various SQL queries to extract actionable business insights like identifying countries with the most rented movies, customers with a penchant for Sci-Fi movies, the most revenue-generating movie, and more.

For this week we will use the movie_rental database. Make sure that you have the [database installed](#) and you are connected to it. You can use [this ERD](#) and [Data Dictionary](#).

Note: All Q1 - Q3 should be done through CLI at the server level (\$) and not through *psql*, including CREATE DATABASE command.

Q1. Using linux specific commands (CLI) ***create a new folder and a SQL backup file of the database movie_rental,*** on your server as ***upxxxxx/db_backup/movie_rental_bk.sql*** . Remember, in order to use Linux specific commands, you need to be at the server level and **not** database (\$).

A:

[EDITABLE CODE]

```
mkdir /home/up2200918/db_backup
cd /home/up2200918/db_backup
pg_dump -U up2200918 movie_rental > movie_rental_bk.sql
```

[OUTPUT SCREENSHOT]

```
up2200918@up2200918:~$ mkdir /home/up2200918/db_backup
up2200918@up2200918:~$ cd /home/up2200918/db_backup
up2200918@up2200918:~/db_backup$ pg_dump -U up2200918 movie_rental > movie_rental_bk.sql
up2200918@up2200918:~/db_backup$ ls
movie_rental_bk.sql
```

Q2. Using linux specific commands (CLI) **delete database movie_rental**.

A:

[EDITABLE CODE]

```
dropdb movie_rental
```

[OUTPUT SCREENSHOT]

```
up2200918@up2200918:~$ dropdb movie_rental
up2200918@up2200918:~$ \l
-bash: 1: command not found
up2200918@up2200918:~$ psql
psql (13.11 (Debian 13.11-0+deb11u1))
Type "help" for help.

up2200918=# \l
```

List of databases					
Name	Owner	Encoding	Collate	Ctype	Access privileges
database_name	up2200918	UTF8	C.UTF-8	C.UTF-8	
dsd_22	postgres	UTF8	C.UTF-8	C.UTF-8	
finding_jane	up2200918	UTF8	C.UTF-8	C.UTF-8	
lab1	up2200918	UTF8	C.UTF-8	C.UTF-8	
lab2	up2200918	UTF8	C.UTF-8	C.UTF-8	
lab3	up2200918	UTF8	C.UTF-8	C.UTF-8	
postgres	postgres	UTF8	C.UTF-8	C.UTF-8	
template0	postgres	UTF8	C.UTF-8	C.UTF-8	=c/postgres + postgres=CTc/postgres
template1	postgres	UTF8	C.UTF-8	C.UTF-8	=c/postgres + postgres=CTc/postgres
test-image-2	test-image-2	UTF8	C.UTF-8	C.UTF-8	
testinglab	up2200918	UTF8	C.UTF-8	C.UTF-8	
uop	uop	UTF8	C.UTF-8	C.UTF-8	
up2200918	up2200918	UTF8	C.UTF-8	C.UTF-8	

```
(13 rows)
```

Q3. Using linux commands, **restore your database from the backup**. If you are encountering issues or for any reason you cannot restore the database from backup, make another clean [installation](#) of the database and start from Q1.

A:

[EDITABLE CODE]

```
-- Create a new clean database
createdb movie_rental
-- Move to backup folder
```

```
cd /home/up2200918/db_backup
-- Restore the database from the backup
psql -U up2200918 movie rental < movie_rental bk.sql
```

[OUTPUT SCREENSHOT]

[illegible]

Q4. Now that the DB is restored, let's see which country has the most rented movies, so the manager will be able to create some marketing strategies. The manager would like to focus on the top 5.

A:

[EDITABLE CODE]

```
SELECT co.country, COUNT(*) as num_rented_movies
FROM customer cu
JOIN rental r ON cu.customer_id = r.customer_id
JOIN address a ON cu.address_id = a.address_id
JOIN city c ON a.city_id = c.city_id
JOIN country co ON c.country_id = co.country_id
GROUP BY co.country
ORDER BY num_rented_movies DESC
```

```
LIMIT 5;
```

[OUTPUT SCREENSHOT]

```
movie_rental=# SELECT co.country, COUNT(*) as num_rented_movies
FROM customer cu
JOIN rental r ON cu.customer_id = r.customer_id
JOIN address a ON cu.address_id = a.address_id
JOIN city c ON a.city_id = c.city_id
JOIN country co ON c.country_id = co.country_id
GROUP BY co.country
ORDER BY num_rented_movies DESC
LIMIT 5;
```

country	num_rented_movies
India	1572
China	1426
United States	968
Japan	825
Mexico	796

(5 rows)

Q5. The manager is a huge Sci-Fi genre fan and would like to find which **customers have rented more than 2 movies** from this genre. Also he would like to run some statistics over the city the customer is from and their contact details (phone and email). Order by city and customer name.

A:

[EDITABLE CODE]

```
SELECT c.first_name, c.last_name, c.email, a.phone, ct.city
FROM customer c
JOIN rental r ON c.customer_id = r.customer_id
JOIN inventory i ON r.inventory_id = i.inventory_id
JOIN movie m ON i.movie_id = m.movie_id
JOIN movie_category mc ON m.movie_id = mc.movie_id
JOIN category cat ON mc.category_id = cat.category_id
JOIN address a ON c.address_id = a.address_id
JOIN city ct ON a.city_id = ct.city_id
WHERE cat.name = 'Sci-Fi'
GROUP BY c.first_name, c.last_name, c.email, a.phone, ct.city
HAVING COUNT(*) > 2
ORDER BY ct.city, c.first_name, c.last_name;
```

[OUTPUT SCREENSHOT]

first_name	last_name	email	phone	city
PEGGY	MYERS	PEGGY.MYERS@sakilacustomer.org	196568435814	Abha
TOM	MILNER	TOM.MILNER@sakilacustomer.org	985109775584	Abu Dhabi
SHERRI	RHODES	SHERRI.RHODES@sakilacustomer.org	510737228015	Ahmadnagar
CLAYTON	BARBEE	CLAYTON.BARBEE@sakilacustomer.org	380077794770	Alvorada
ARLENE	HARVEY	ARLENE.HARVEY@sakilacustomer.org	460795526514	Ambattur
JANET	PHILLIPS	JANET.PHILLIPS@sakilacustomer.org	675292816413	Antofagasta
CRYSTAL	FORD	CRYSTAL.FORD@sakilacustomer.org	709935135487	Ashdod
MAUREEN	LITTLE	MAUREEN.LITTLE@sakilacustomer.org	684192903087	Asuncin
LINDA	WILLIAMS	LINDA.WILLIAMS@sakilacustomer.org	448477190408	Athenai
JIMMY	SCHRADER	JIMMY.SCHRADER@sakilacustomer.org	604177838256	Atinsk
TED	BREAUX	TED.BREAUX@sakilacustomer.org	488600270038	Baicheng
ANDREW	PURDY	ANDREW.PURDY@sakilacustomer.org	119501405123	Baku
BRANDON	HUEY	BRANDON.HUEY@sakilacustomer.org	99883471275	Balikesir
MARIO	CHEATHAM	MARIO.CHEATHAM@sakilacustomer.org	406784385440	Batna
ELMER	NOE	ELMER.NOE@sakilacustomer.org	448876499197	Battambang
JUDY	GRAY	JUDY.GRAY@sakilacustomer.org	107137400143	Bchar
MARSHA	DOUGLAS	MARSHA.DOUGLAS@sakilacustomer.org	245477603573	Beira
WILMA	RICHARDS	WILMA.RICHARDS@sakilacustomer.org	168758068397	Bellevue
CARLOS	COUGHLIN	CARLOS.COUGHLIN@sakilacustomer.org	924815207181	Bhavnagar
ERIK	GUILLEN	ERIK.GUILLEN@sakilacustomer.org	80593242951	Bhimavaram
TIM	CARY	TIM.CARY@sakilacustomer.org	195337700615	Bijapur
JESSIE	MILAM	JESSIE.MILAM@sakilacustomer.org	383353187467	Binzhou
KEVIN	SCHULER	KEVIN.SCHULER@sakilacustomer.org	908029859266	Birgunj
JOY	GEORGE	JOY.GEORGE@sakilacustomer.org	479007229460	Botosani
ANNE	POWELL	ANNE.POWELL@sakilacustomer.org	720998247660	Bradford
MARC	OUTLAW	MARC.OUTLAW@sakilacustomer.org	465897838272	Brindisi
VALERIE	BLACK	VALERIE.BLACK@sakilacustomer.org	885899703621	Brookton
BILLY	POULIN	BILLY.POULIN@sakilacustomer.org	333390595558	Cabuyao
DAVID	ROYAL	DAVID.ROYAL@sakilacustomer.org	504434452842	Callao
WILLARD	LUMPKIN	WILLARD.LUMPKIN@sakilacustomer.org	377633994405	Carmen
EVA	RAMOS	EVA.RAMOS@sakilacustomer.org	9099941466	Clarksville
STELLA	MORENO	STELLA.MORENO@sakilacustomer.org	266798132374	Coacalco de Berriozbal
MAURICE	CRAWLEY	MAURICE.CRAWLEY@sakilacustomer.org	684529463244	Coatzacoalcas
RENEE	LANE	RENEE.LANE@sakilacustomer.org	662227486184	Compton
JEFFERY	PINSON	JEFFERY.PINSON@sakilacustomer.org	15273765306	Dadu
BRYAN	HARDISON	BRYAN.HARDISON@sakilacustomer.org	775235029633	Dallas
MAE	FLETCHER	MAE.FLETCHER@sakilacustomer.org	96604821070	Donostia-San Sebastin
VELMA	LUCAS	VELMA.LUCAS@sakilacustomer.org	558236142492	Effon-Alaiye
JON	WILES	JON.WILES@sakilacustomer.org	205524798287	El Alto
JIM	REA	JIM.REA@sakilacustomer.org	524567129902	El Fuerte
DUSTIN	GILLETTE	DUSTIN.GILLETTE@sakilacustomer.org	131912793873	Emmen
MICHAEL	FORMAN	MICHAEL.FORMAN@sakilacustomer.org	411549550611	Escobar
EDGAR	RHOADS	EDGAR.RHOADS@sakilacustomer.org	402630109080	Eskisehir
MORRIS	MCCARTER	MORRIS.MCCARTER@sakilacustomer.org	278669994384	Fengshan
ALICE	STEWART	ALICE.STEWART@sakilacustomer.org	171822533480	Fontana
DEBBIE	REYES	DEBBIE.REYES@sakilacustomer.org	581852137991	Fukuyama
(152 rows)				
RONALD	WEINER	RONALD.WEINER@sakilacustomer.org	674805712553	San Felipe del Progreso
CARL	ARTIS	CARL.ARTIS@sakilacustomer.org	20064292617	San Lorenzo
MARION	SNYDER	MARION.SNYDER@sakilacustomer.org	391065549876	Santa Brbara dOeste
ERIC	ROBERT	ERIC.ROBERT@sakilacustomer.org	105470691550	Santa F
JUSTIN	NGO	JUSTIN.NGO@sakilacustomer.org	764680915323	Santo Andr
DEBORAH	WALKER	DEBORAH.WALKER@sakilacustomer.org	196495945706	Shikarpur
SHERRY	MARSHALL	SHERRY.MARSHALL@sakilacustomer.org	787654415858	Shubra al-Khayma
KATHRYN	COLEMAN	KATHRYN.COLEMAN@sakilacustomer.org	821972242086	Simferopol
WANDA	PATTERSON	WANDA.PATTERSON@sakilacustomer.org	198123170793	Sincelejo
JUNE	CARROLL	JUNE.CARROLL@sakilacustomer.org	506134035434	Skikda
LUCILLE	HOLMES	LUCILLE.HOLMES@sakilacustomer.org	918032330119	Soshanguve
GREGORY	MAULDIN	GREGORY.MAULDIN@sakilacustomer.org	80303246192	Sousse
BOBBY	BOUDREAU	BOBBY.BOUDREAU@sakilacustomer.org	934352415130	South Hill
LOUISE	JENKINS	LOUISE.JENKINS@sakilacustomer.org	800716535041	Springs
FREDDIE	DUGGAN	FREDDIE.DUGGAN@sakilacustomer.org	644021380889	Sullana
CHARLES	KOWALSKI	CHARLES.KOWALSKI@sakilacustomer.org	181179321332	Sungai Petani
JAY	ROBB	JAY.ROBB@sakilacustomer.org	834061016202	Surakarta
TERRANCE	ROUSH	TERRANCE.ROUSH@sakilacustomer.org	437829801725	Szkesfehrvr
DIANNE	SHELTON	DIANNE.SHELTON@sakilacustomer.org	117592274996	Tabriz
CAROLINE	BOWMAN	CAROLINE.BOWMAN@sakilacustomer.org	435785045362	Tallahassee
PHILLIP	HOLM	PHILLIP.HOLM@sakilacustomer.org	776031833752	Tama
BECKY	MILES	BECKY.MILES@sakilacustomer.org	648482415405	Tambaram
MICHELLE	CLARK	MICHELLE.CLARK@sakilacustomer.org	892775750063	Tangail
CLIFTON	MALCOLM	CLIFTON.MALCOLM@sakilacustomer.org	29341849811	Tanshui
AUSTIN	CINTRON	AUSTIN.CINTRON@sakilacustomer.org	288241215394	Tieli
CASEY	MENA	CASEY.MENA@sakilacustomer.org	525518075499	Tokat
PAULINE	HENRY	PAULINE.HENRY@sakilacustomer.org	321944036800	Torren
VIRGINIA	GREEN	VIRGINIA.GREEN@sakilacustomer.org	440512153169	Toulouse
EDWARD	BAUGH	EDWARD.BAUGH@sakilacustomer.org	46568045367	Trshavn
RUSSELL	BRINSON	RUSSELL.BRINSON@sakilacustomer.org	821476736117	Tychy
BEATRICE	ARNOLD	BEATRICE.ARNOLD@sakilacustomer.org	264541743403	Udaipur
ANGEL	BARCLAY	ANGEL.BARCLAY@sakilacustomer.org	759586584889	Ueda
JOANNE	ROBERTSON	JOANNE.ROBERTSON@sakilacustomer.org	597815221267	Urawa
EVELYN	MORGAN	EVELYN.MORGAN@sakilacustomer.org	889318963672	Vaduz
KRISTINA	CHAMBERS	KRISTINA.CHAMBERS@sakilacustomer.org	892523334	Valle de la Pascua
MIKE	WAY	MIKE.WAY@sakilacustomer.org	206169448769	Valparai
WILLIE	HOWELL	WILLIE.HOWELL@sakilacustomer.org	991802825778	Vicente Lpez
STEPHANIE	MITCHELL	STEPHANIE.MITCHELL@sakilacustomer.org	42384721397	Yerevan
GILBERT	SLEDGE	GILBERT.SLEDGE@sakilacustomer.org	959467760895	York
MICHELE	GRANT	MICHELE.GRANT@sakilacustomer.org	499408708580	Yuncheng
CONSTANCE	REID	CONSTANCE.REID@sakilacustomer.org	588964509072	Zaria
RONNIE	RICKETTS	RONNIE.RICKETTS@sakilacustomer.org	670370974122	Ziguinchor
SETH	HANNON	SETH.HANNON@sakilacustomer.org	864392582257	al-Manama
JOHNNY	TURPIN	JOHNNY.TURPIN@sakilacustomer.org	765957414528	al-Qadarif
TRAVIS	ESTEP	TRAVIS.ESTEP@sakilacustomer.org	214976066017	al-Qatif
CHARLOTTE	HUNTER	CHARLOTTE.HUNTER@sakilacustomer.org	935448624185	guas Lindas de Gois
(END)				

Q6. The business is on the right track, however it seems that in May a particular movie was quite popular generating a high revenue. What movie was and what was the amount generated?

A:

[EDITABLE CODE]

```
SELECT m.title as "Movie Title",
SUM(all_payments.amount) AS "Total Revenue"
FROM
(
    SELECT rental_id, amount, payment_date FROM payment_p2022_01
    UNION ALL
    SELECT rental_id, amount, payment_date FROM payment_p2022_02
    UNION ALL
    SELECT rental_id, amount, payment_date FROM payment_p2022_03
    -- Add more UNION ALL statements for each month up to payment_p2022_07
    UNION ALL
    SELECT rental_id, amount, payment_date FROM payment_p2022_04
    UNION ALL
    SELECT rental_id, amount, payment_date FROM payment_p2022_05
    UNION ALL
    SELECT rental_id, amount, payment_date FROM payment_p2022_06
    UNION ALL
    SELECT rental_id, amount, payment_date FROM payment_p2022_07
) AS all_payments
JOIN rental r ON all_payments.rental_id = r.rental_id
JOIN inventory i ON r.inventory_id = i.inventory_id
JOIN movie m ON i.movie_id = m.movie_id
WHERE date_part('month', all_payments.payment_date) = 5 -- 5 = May Month
GROUP BY m.title
ORDER BY "Total Revenue" DESC
LIMIT 1; -- Top Movie Revenue for May
```


[OUTPUT SCREENSHOT]

```

movie_rental=# SELECT m.title as "Movie Title",
SUM(all_payments.amount) AS "Total Revenue"
FROM
(
  SELECT rental_id, amount, payment_date FROM payment_p2022_01
  UNION ALL
  SELECT rental_id, amount, payment_date FROM payment_p2022_02
  UNION ALL
  SELECT rental_id, amount, payment_date FROM payment_p2022_03
  -- Add more UNION ALL statements for each month up to payment_p2022_07
  UNION ALL
  SELECT rental_id, amount, payment_date FROM payment_p2022_04
  UNION ALL
  SELECT rental_id, amount, payment_date FROM payment_p2022_05
  UNION ALL
  SELECT rental_id, amount, payment_date FROM payment_p2022_06
  UNION ALL
  SELECT rental_id, amount, payment_date FROM payment_p2022_07
) AS all_payments
JOIN rental r ON all_payments.rental_id = r.rental_id
JOIN inventory i ON r.inventory_id = i.inventory_id
JOIN movie m ON i.movie_id = m.movie_id
WHERE date_part('month', all_payments.payment_date) = 5 -- 5 = May Month
GROUP BY m.title
ORDER BY "Total Revenue" DESC
LIMIT 1; -- Top Movie Revenue for May
 Movie Title | Total Revenue
-----+-----
 DOGMA FAMILY |          74.89
(1 row)

```

Q7. Now the manager wants to have a closer look at the general business income. He would like to see the **top 5 categories total gross revenue for January per category**.

A:

[EDITABLE CODE]

```

SELECT
  cat.name AS category_name,
  SUM(all_payments.amount) AS total_revenue
FROM
(
  SELECT rental_id, amount, payment_date FROM payment_p2022_01
  UNION ALL
  SELECT rental_id, amount, payment_date FROM payment_p2022_02
  UNION ALL
  SELECT rental_id, amount, payment_date FROM payment_p2022_03
  -- Add more UNION ALL statements for each month up to payment_p2022_07
  UNION ALL
  SELECT rental_id, amount, payment_date FROM payment_p2022_04
  UNION ALL
  SELECT rental_id, amount, payment_date FROM payment_p2022_05
  UNION ALL
  SELECT rental_id, amount, payment_date FROM payment_p2022_06
  UNION ALL
  SELECT rental_id, amount, payment_date FROM payment_p2022_07
) AS all_payments
JOIN rental r ON all_payments.rental_id = r.rental_id
JOIN inventory i ON r.inventory_id = i.inventory_id
JOIN movie m ON i.movie_id = m.movie_id
JOIN movie_category mc ON m.movie_id = mc.movie_id
JOIN category cat ON mc.category_id = cat.category_id

```

```
WHERE date_part('month', all_payments.payment_date) = 1
GROUP BY cat.name
ORDER BY total_revenue DESC
LIMIT 5;
```

[OUTPUT SCREENSHOT]

```
movie_rental=# SELECT
  cat.name AS category_name,
  SUM(all_payments.amount) AS total_revenue
FROM
  (
    SELECT rental_id, amount, payment_date FROM payment_p2022_01
    UNION ALL
    SELECT rental_id, amount, payment_date FROM payment_p2022_02
    UNION ALL
    SELECT rental_id, amount, payment_date FROM payment_p2022_03
    -- Add more UNION ALL statements for each month up to payment_p2022_07
    UNION ALL
    SELECT rental_id, amount, payment_date FROM payment_p2022_04
    UNION ALL
    SELECT rental_id, amount, payment_date FROM payment_p2022_05
    UNION ALL
    SELECT rental_id, amount, payment_date FROM payment_p2022_06
    UNION ALL
    SELECT rental_id, amount, payment_date FROM payment_p2022_07
  ) AS all_payments
JOIN rental r ON all_payments.rental_id = r.rental_id
JOIN inventory i ON r.inventory_id = i.inventory_id
JOIN movie m ON i.movie_id = m.movie_id
JOIN movie_category mc ON m.movie_id = mc.movie_id
JOIN category cat ON mc.category_id = cat.category_id
WHERE date_part('month', all_payments.payment_date) = 1
GROUP BY cat.name
ORDER BY total_revenue DESC
LIMIT 5;
```

category_name	total_revenue
Sci-Fi	248.41
Action	247.42
New	237.46
Sports	235.48
Comedy	231.56

(5 rows)

- Q8.** The *Bottleneck* concept refers to:
- A system that will allow easy query
 - Aspects that will limit database performance
 - Display critical business metrics and gather reports in one place

DBPRIN – UP2200918

- d. The absolute minimum interruption to the process
- e. Process mapping the workflow visually to spot congestions

Q9. Performance of a databases is measured against:

- a. Queries
- b. Business rules
- c. Software requirements
- d. Benchmarks**
- e. Monitoring tools

Q10. How many tuning levels are for databases?

- a. 6
- b. 4
- c. 2
- d. 3**
- e. 1

Conclusion/Reflection

Lab 10 was an incredibly enriching exercise that solidified my understanding of database administration and its intersection with Linux command-line operations. Creating a database backup, deleting the database, and then restoring it from the backup file illustrated the importance of backup systems in real-world scenarios.

The lab also helped me appreciate how powerful SQL can be when it comes to extracting valuable business insights. By executing specific queries, I could determine patterns and trends that could potentially be leveraged for business decision-making and strategy planning.

Another important aspect that Lab 10 emphasized was the concept of database performance and bottleneck identification. The ability to measure and optimize database performance is a vital skill in the field of data management.

Overall, Lab 10 was a challenging but enlightening experience, effectively honing my database administration skills, enhancing my SQL expertise, and broadening my understanding of real-world database use-cases. This knowledge and these skills will undoubtedly prove to be crucial in my future data-related endeavors.

LAB 11 (26/07/2023) - Christmas Fun

Lab Description

Lab 11 was an engaging exploration into advanced PostgreSQL database management and SQL queries. It covered a wide range of topics including recursive CTEs, joins, subqueries, full-text search, and query optimization. Implemented these techniques on the Movie Rental database, which provided a practical and enriching environment for improving query writing and data manipulation skills.

In addition to these advanced techniques, also introduced the fascinating world of PostgreSQL tricks such as generating automatic number series, creating festive graphical outputs like Christmas trees and forests, generating random numbers, and simulating UUIDs.

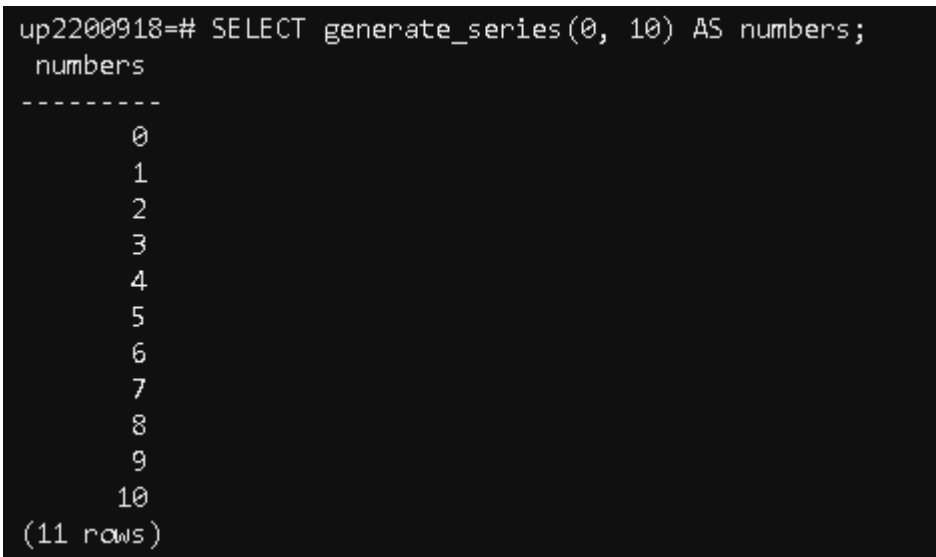
Q1. Create an automatic series of numbers from 0 to 10. The output should be generated by PgSQL without any database, table or insert. No specific database is needed.

A:

[EDITABLE CODE]


```
SELECT generate_series(0, 10) AS numbers;
```

[OUTPUT SCREENSHOT]



```
up2200918=# SELECT generate_series(0, 10) AS numbers;
 numbers
-----
      0
      1
      2
      3
      4
      5
      6
      7
      8
      9
     10
(11 rows)
```

Q2. We want to get festive with our PostgreSQL. Can you generate a Christmas Tree like below? No specific database, table or insert is required.



A:

[EDITABLE CODE]

```
WITH RECURSIVE small_tree(tree_depth, "Christmas Tree") AS (
    SELECT 1 AS tree_depth, rpad(' ', 10, ' ') || '*' AS "Christmas Tree"
    FROM generate_series(1, 1)
    UNION
    SELECT small_tree.tree_depth + 1 AS tree_depth,
           rpad(' ', 10 - small_tree.tree_depth, ' ') ||
           rpad('*', small_tree.tree_depth + 1, '.') ||
           lpad('*', small_tree.tree_depth, '.') AS "Christmas Tree"
    FROM small_tree
    WHERE small_tree.tree_depth < 10
)

SELECT "Christmas Tree"
FROM small_tree;
```

[OUTPUT SCREENSHOT]

```

up2200918=# WITH RECURSIVE small_tree(tree_depth, "Christmas Tree") AS (
    SELECT 1 AS tree_depth, rpad(' ', 10, '*') || '*' AS "Christmas Tree"
    FROM generate_series(1, 1)
    UNION
    SELECT small_tree.tree_depth + 1 AS tree_depth,
           rpad(' ', 10 - small_tree.tree_depth, '*') ||
           rpad('*', small_tree.tree_depth + 1, '.') ||
           lpad('*', small_tree.tree_depth, '.') AS "Christmas Tree"
    FROM small_tree
    WHERE small_tree.tree_depth < 10
)
SELECT "Christmas Tree"
FROM small_tree;
      Christmas Tree
-----
          *
        *. *
       *. . *.
      *. . . *.
     *. . . . *.
    *. . . . . *.
   *. . . . . *.
  *. . . . . *.
 *. . . . . *.
*. . . . . *.
*. . . . . *.
(10 rows)

```

Q3. How about a small forest of them?



A:

[EDITABLE CODE]

```

WITH RECURSIVE small_tree(tree_depth, "Christmas Tree") AS (
    SELECT 1 AS tree_depth, rpad(' ', 10, ' ') || '*' AS "Christmas Tree"
    FROM generate_series(1, 1)
    UNION
    SELECT small_tree.tree_depth + 1 AS tree_depth,
           rpad(' ', 10 - small_tree.tree_depth, ' ') ||
           rpad('*', small_tree.tree_depth + 1, '.') ||
           lpad('*', small_tree.tree_depth, '.') AS "Christmas Tree"
    FROM small_tree
    WHERE small_tree.tree_depth < 10
)
SELECT
    t1."Christmas Tree" AS "Christmas Tree",
    t2."Christmas Tree" AS "Christmas Tree",
    t3."Christmas Tree" AS "Christmas Tree"
FROM
    small_tree t1
JOIN
    small_tree t2 ON t2.tree_depth = t1.tree_depth
JOIN
    small_tree t3 ON t3.tree_depth = t1.tree_depth
WHERE
    t1.tree_depth <= 10;

```

[OUTPUT SCREENSHOT]

```

up2200918=# WITH RECURSIVE small_tree(tree_depth, "Christmas Tree") AS (
  SELECT 1 AS tree_depth, rpad(' ', 10, ' ') || '*' AS "Christmas Tree"
  FROM generate_series(1, 1)
  UNION
  SELECT small_tree.tree_depth + 1 AS tree_depth,
         rpad(' ', 10 - small_tree.tree_depth, ' ') ||
         rpad('*', small_tree.tree_depth + 1, ' ') ||
         lpad('*', small_tree.tree_depth, ' ') AS "Christmas Tree"
  FROM small_tree
  WHERE small_tree.tree_depth < 10
)
SELECT
  t1."Christmas Tree" AS "Christmas Tree",
  t2."Christmas Tree" AS "Christmas Tree",
  t3."Christmas Tree" AS "Christmas Tree"
FROM
  small_tree t1
JOIN
  small_tree t2 ON t2.tree_depth = t1.tree_depth
JOIN
  small_tree t3 ON t3.tree_depth = t1.tree_depth
WHERE
  t1.tree_depth <= 10;

```

Christmas Tree	Christmas Tree	Christmas Tree
<pre> * * _ * * _ _ _ * * _ _ _ _ _ * * _ _ _ _ _ * * _ _ _ _ _ * * _ _ _ _ _ * * _ _ _ _ _ * * _ _ _ _ _ * * _ _ _ _ _ * </pre>	<pre> * * _ * * _ _ _ * * _ _ _ _ _ * * _ _ _ _ _ * * _ _ _ _ _ * * _ _ _ _ _ * * _ _ _ _ _ * * _ _ _ _ _ * * _ _ _ _ _ * </pre>	<pre> * * _ * * _ _ _ * * _ _ _ _ _ * * _ _ _ _ _ * * _ _ _ _ _ * * _ _ _ _ _ * * _ _ _ _ _ * * _ _ _ _ _ * * _ _ _ _ _ * </pre>

(10 rows)

Q4. Put all your SQL knowledge at work and based on the Movie Rental database, what is the best query you can create?

A:

[EDITABLE CODE]

```

-- Top 5 most popular movie categories (genres) based on
the number of rentals in the past year:
SELECT
  cat.name AS category_name,
  COUNT(*) AS rental_count
FROM
  rental r
JOIN inventory i ON r.inventory_id = i.inventory_id
JOIN movie m ON i.movie_id = m.movie_id
JOIN movie_category mc ON m.movie_id = mc.movie_id
JOIN category cat ON mc.category_id = cat.category_id
WHERE
  r.rental_date >= CURRENT_DATE - INTERVAL '1 year'
GROUP BY
  cat.name
ORDER BY
  rental_count DESC
LIMIT 5;

```

[OUTPUT SCREENSHOT]

```

category_name | rental_count
-----+-----
Sports        |           672
Animation     |           645
Sci-Fi        |           612
Family        |           606
Action        |           604
(5 rows)

```

Q5. What “cool tricks” can you show with PostgreSQL? Similar to Q1 - Q3?

A:

[EDITABLE CODE]

```

--Generate a Random Number: return a random number between 0 and 1.
SELECT RANDOM() AS random_number;

--Simulate UUID Generation: This query will generate a simulated UUID.
SELECT md5(random()::text || clock_timestamp()::text)::uuid AS
simulated_uuid;

--Date Arithmetic: Perform date arithmetic in PostgreSQL to add or subtract
days, months, or years to/from a date.
SELECT CURRENT_DATE AS current_date,
CURRENT_DATE + INTERVAL '1 day' AS next_day,
CURRENT_DATE - INTERVAL '1 month' AS last_month,
CURRENT_DATE + INTERVAL '1 year' AS next_year;

```

[OUTPUT SCREENSHOT]

```

up2200918=# --Generate a Random Number: return a random number between 0 and 1.
SELECT RANDOM() AS random_number;

--Simulate UUID Generation: This query will generate a simulated UUID.
SELECT md5(random()::text || clock_timestamp()::text)::uuid AS simulated_uuid;

--Date Arithmetic: Perform date arithmetic in PostgreSQL to add or subtract days, months, or years to/from a date.
SELECT CURRENT_DATE AS current_date,
CURRENT_DATE + INTERVAL '1 day' AS next_day,
CURRENT_DATE - INTERVAL '1 month' AS last_month,
CURRENT_DATE + INTERVAL '1 year' AS next_year;
 random_number
-----
0.9609438909779229
(1 row)

 simulated_uuid
-----
82c7fa73-c707-a45a-d334-9ac960005228
(1 row)

 current_date | next_day | last_month | next_year
-----+-----+-----+-----
2023-07-21 | 2023-07-22 00:00:00 | 2023-06-21 00:00:00 | 2024-07-21 00:00:00
(1 row)

```


Conclusion/Reflection

Lab 11 was a valuable experience that deepened my understanding of PostgreSQL and advanced SQL techniques. I found the topic of recursive CTEs intriguing and the opportunity to use this to generate graphical outputs like Christmas trees was fun and enlightening. It was fascinating to see how PostgreSQL can be used creatively, beyond the traditional database operations.

Working on complex queries with the Movie Rental database also allowed me to appreciate the importance of query optimization in real-world scenarios. Extracting the most popular movie categories based on the number of rentals in the past year showed how SQL can provide valuable business insights.

The "cool tricks" with PostgreSQL was another highlight of Lab 11. Generating random numbers and simulating UUIDs are techniques that I can see being quite handy in many real-world applications.

Overall, Lab 11 was a rewarding and enriching experience that improved my SQL skills, broadened my understanding of advanced database management, and exposed me to the unique capabilities of PostgreSQL. I feel more confident in handling complex database operations and optimizing performance for real-world database scenarios.